

Final_Project_Code

December 5, 2022

1 Import the Dataset

```
[1]: import pandas as pd
df=pd.read_csv("https://raw.githubusercontent.com/MaramShriem/
↳-Marketing-Dataset/main/bank-full.csv",sep=';')
df.to_csv(r'TDMaketing.csv', index = False)#to save the dataset as CSV
```

2 EDA

Dataset Description

```
[2]: df.describe(include='all')
```

```
[2]:
```

	age	job	marital	education	default	balance \
count	45211.000000	45211	45211	45211	45211	45211.000000
unique	NaN	12	3	4	2	NaN
top	NaN	blue-collar	married	secondary	no	NaN
freq	NaN	9732	27214	23202	44396	NaN
mean	40.936210	NaN	NaN	NaN	NaN	1362.272058
std	10.618762	NaN	NaN	NaN	NaN	3044.765829
min	18.000000	NaN	NaN	NaN	NaN	-8019.000000
25%	33.000000	NaN	NaN	NaN	NaN	72.000000
50%	39.000000	NaN	NaN	NaN	NaN	448.000000
75%	48.000000	NaN	NaN	NaN	NaN	1428.000000
max	95.000000	NaN	NaN	NaN	NaN	102127.000000

	housing	loan	contact	day	month	duration \
count	45211	45211	45211	45211.000000	45211	45211.000000
unique	2	2	3	NaN	12	NaN
top	yes	no	cellular	NaN	may	NaN
freq	25130	37967	29285	NaN	13766	NaN
mean	NaN	NaN	NaN	15.806419	NaN	258.163080
std	NaN	NaN	NaN	8.322476	NaN	257.527812
min	NaN	NaN	NaN	1.000000	NaN	0.000000
25%	NaN	NaN	NaN	8.000000	NaN	103.000000

50%	NaN	NaN	NaN	16.000000	NaN	180.000000
75%	NaN	NaN	NaN	21.000000	NaN	319.000000
max	NaN	NaN	NaN	31.000000	NaN	4918.000000

	campaign	pdays	previous	poutcome	y
count	45211.000000	45211.000000	45211.000000	45211	45211
unique	NaN	NaN	NaN	4	2
top	NaN	NaN	NaN	unknown	no
freq	NaN	NaN	NaN	36959	39922
mean	2.763841	40.197828	0.580323	NaN	NaN
std	3.098021	100.128746	2.303441	NaN	NaN
min	1.000000	-1.000000	0.000000	NaN	NaN
25%	1.000000	-1.000000	0.000000	NaN	NaN
50%	2.000000	-1.000000	0.000000	NaN	NaN
75%	3.000000	-1.000000	0.000000	NaN	NaN
max	63.000000	871.000000	275.000000	NaN	NaN

```
[3]: df.shape[1]
```

```
[3]: 17
```

```
[4]: df.head()
```

```
[4]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no

Check the datatypes of the attributes

```
[5]: print(df.dtypes)
```

```
age          int64
job          object
marital      object
education    object
default      object
balance      int64
housing      object
```

```

loan          object
contact       object
day           int64
month         object
duration      int64
campaign      int64
pdays        int64
previous      int64
poutcome     object
y            object
dtype: object

```

Checking the missing values

```

[6]: df.isnull().sum()
     #the answer is No

```

```

[6]: age          0
     job          0
     marital      0
     education    0
     default      0
     balance      0
     housing      0
     loan         0
     contact      0
     day          0
     month        0
     duration     0
     campaign     0
     pdays        0
     previous     0
     poutcome     0
     y            0
     dtype: int64

```

Check if the data is balanced or Imbalanced

```

[7]: x=df.drop(["y"],axis=1)
     y=df["y"]
     print(y.value_counts())

```

```

no      39922
yes      5289
Name: y, dtype: int64


```

```

[8]: y.value_counts().plot.pie(autopct="%0.2f")
     #we have only 11.7% of the yes class in the dataset; that means the data is
     ↳ Imbalanced.

```

```
[8]: <AxesSubplot: ylabel='y'>
```



Plot the correlation between numerical variables

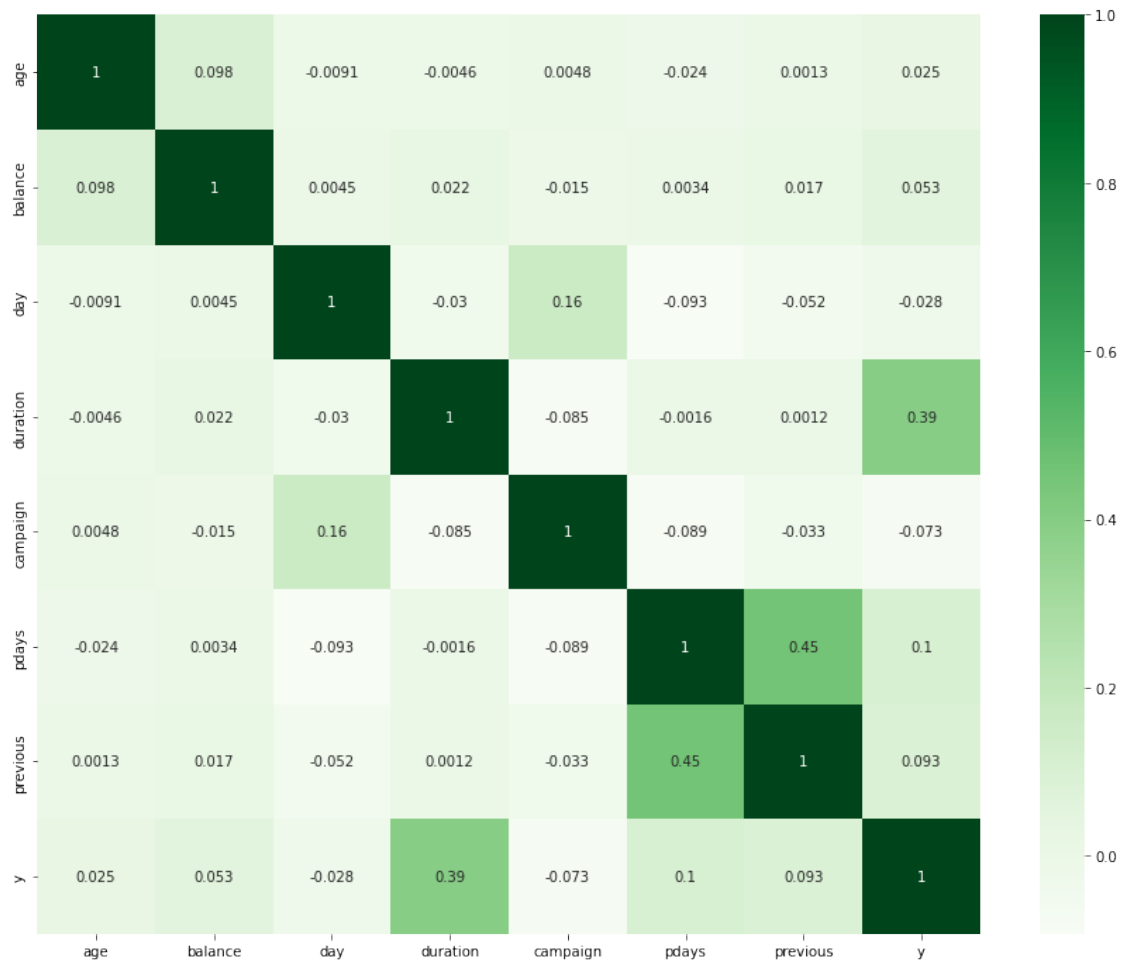
```
[9]: df['y'].replace({'no':0, 'yes':1}, inplace=True)
df.corr()
```

```
[9]:
```

	age	balance	day	duration	campaign	pdays	\
age	1.000000	0.097783	-0.009120	-0.004648	0.004760	-0.023758	
balance	0.097783	1.000000	0.004503	0.021560	-0.014578	0.003435	
day	-0.009120	0.004503	1.000000	-0.030206	0.162490	-0.093044	
duration	-0.004648	0.021560	-0.030206	1.000000	-0.084570	-0.001565	
campaign	0.004760	-0.014578	0.162490	-0.084570	1.000000	-0.088628	
pdays	-0.023758	0.003435	-0.093044	-0.001565	-0.088628	1.000000	
previous	0.001288	0.016674	-0.051710	0.001203	-0.032855	0.454820	
y	0.025155	0.052838	-0.028348	0.394521	-0.073172	0.103621	

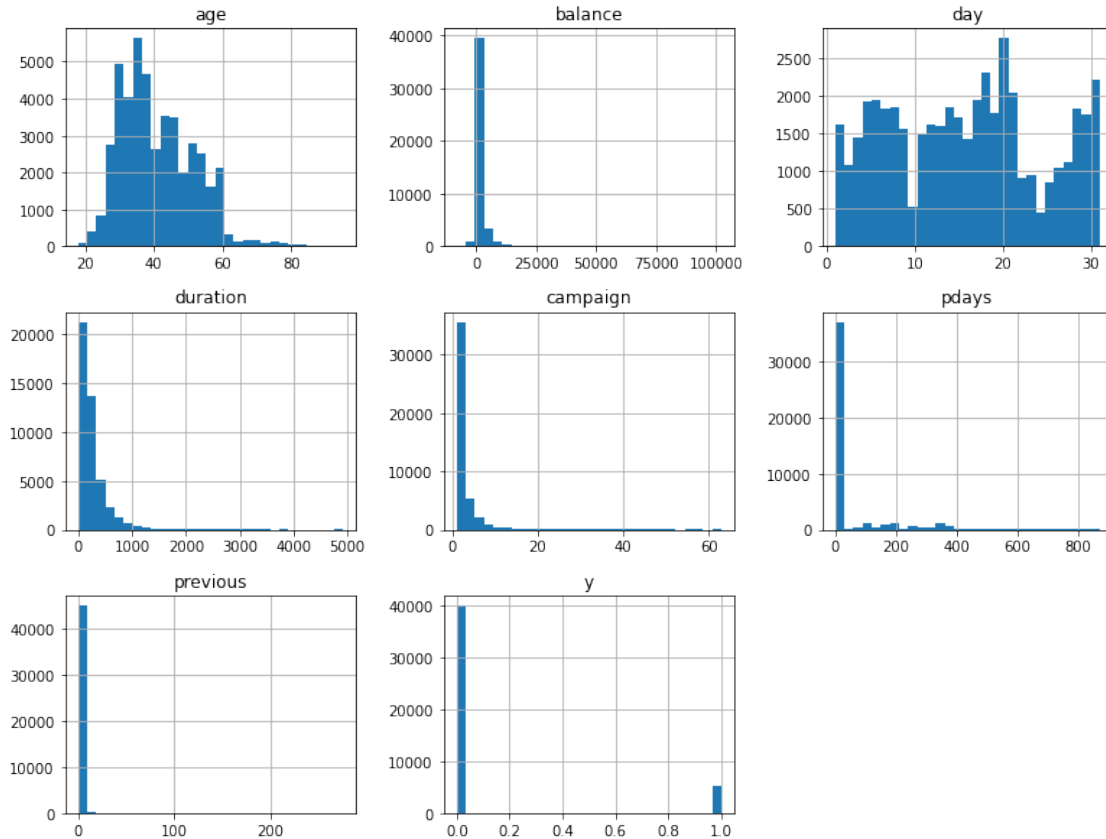
	previous	y
age	0.001288	0.025155
balance	0.016674	0.052838
day	-0.051710	-0.028348
duration	0.001203	0.394521
campaign	-0.032855	-0.073172
pdays	0.454820	0.103621
previous	1.000000	0.093236
y	0.093236	1.000000

```
[10]: import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize = (15, 12))
cor = df.corr()
sns.heatmap(cor, annot = True, cmap = plt.cm.Greens)
plt.show()
```



Histograms for numerical variables

```
[11]: axList = df.hist(bins=29,figsize = (13, 10))
plt.savefig("Hist.png")
```



Correlation for categorical data

[12]: `!pip install dython`

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Requirement already satisfied: dython in /usr/local/lib/python3.8/dist-packages
(0.7.2)
Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.8/dist-
packages (from dython) (0.11.2)
Requirement already satisfied: scikit-plot>=0.3.7 in
/usr/local/lib/python3.8/dist-packages (from dython) (0.3.7)
Requirement already satisfied: scipy>=1.7.1 in /usr/local/lib/python3.8/dist-
packages (from dython) (1.7.3)
Requirement already satisfied: psutil>=5.9.1 in /usr/local/lib/python3.8/dist-
packages (from dython) (5.9.4)
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.8/dist-
packages (from dython) (1.21.6)
Requirement already satisfied: matplotlib>=3.4.3 in
/usr/local/lib/python3.8/dist-packages (from dython) (3.6.2)
Requirement already satisfied: scikit-learn>=0.24.2 in
```

```

/usr/local/lib/python3.8/dist-packages (from dython) (1.1.3)
Requirement already satisfied: pandas>=1.3.2 in /usr/local/lib/python3.8/dist-
packages (from dython) (1.3.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-
packages (from matplotlib>=3.4.3->dython) (21.3)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (1.4.4)
Requirement already satisfied: pyparsing>=2.2.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (3.0.9)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.8/dist-
packages (from matplotlib>=3.4.3->dython) (7.1.2)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (1.0.6)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (2.8.2)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.8/dist-
packages (from matplotlib>=3.4.3->dython) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (4.38.0)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-
packages (from pandas>=1.3.2->dython) (2022.6)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-
packages (from python-dateutil>=2.7->matplotlib>=3.4.3->dython) (1.15.0)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.8/dist-
packages (from scikit-learn>=0.24.2->dython) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.24.2->dython)
(3.1.0)

```

```
[13]: ! pip install -U Imbalanced-learn
```

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Requirement already satisfied: Imbalanced-learn in
/usr/local/lib/python3.8/dist-packages (0.9.1)
Requirement already satisfied: scikit-learn>=1.1.0 in
/usr/local/lib/python3.8/dist-packages (from Imbalanced-learn) (1.1.3)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.8/dist-
packages (from Imbalanced-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.8/dist-packages (from Imbalanced-learn) (3.1.0)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-
packages (from Imbalanced-learn) (1.7.3)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.8/dist-
packages (from Imbalanced-learn) (1.21.6)

```

```
[14]: !pip install matplotlib>=3.4.3
!pip install dython
import pandas as pd
df=pd.read_csv("https://raw.githubusercontent.com/MaramShriem/
↳-Marketing-Dataset/main/bank-full.csv",sep=';')
df.to_csv(r'TDMaketing.csv', index = False)

from dython.nominal import associations
selected_column=
↳df[["job","housing","marital","loan","default","education","contact","month","outcome","y"]
categorical_df = selected_column.copy()
categorical_correlation= associations(categorical_df, filename=
↳'categorical_correlation.png', figsize=(12,12))
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Requirement already satisfied: dython in /usr/local/lib/python3.8/dist-packages (0.7.2)

Requirement already satisfied: matplotlib>=3.4.3 in /usr/local/lib/python3.8/dist-packages (from dython) (3.6.2)

Requirement already satisfied: psutil>=5.9.1 in /usr/local/lib/python3.8/dist-packages (from dython) (5.9.4)

Requirement already satisfied: scikit-learn>=0.24.2 in /usr/local/lib/python3.8/dist-packages (from dython) (1.1.3)

Requirement already satisfied: pandas>=1.3.2 in /usr/local/lib/python3.8/dist-packages (from dython) (1.3.5)

Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.8/dist-packages (from dython) (0.11.2)

Requirement already satisfied: scipy>=1.7.1 in /usr/local/lib/python3.8/dist-packages (from dython) (1.7.3)

Requirement already satisfied: scikit-plot>=0.3.7 in /usr/local/lib/python3.8/dist-packages (from dython) (0.3.7)

Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.8/dist-packages (from dython) (1.21.6)

Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (7.1.2)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (4.38.0)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (21.3)

Requirement already satisfied: cyclo>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (0.11.0)

Requirement already satisfied: pyparsing>=2.2.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (3.0.9)


Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (1.0.6)

Requirement already satisfied: kiwisolver>=1.0.1 in


```

/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (1.4.4)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.8/dist-packages (from matplotlib>=3.4.3->dython) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-
packages (from pandas>=1.3.2->dython) (2022.6)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-
packages (from python-dateutil>=2.7->matplotlib>=3.4.3->dython) (1.15.0)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.8/dist-
packages (from scikit-learn>=0.24.2->dython) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.24.2->dython)
(3.1.0)

```



3 Data Preprocessing:

```
[15]: dff=df.copy()
```

```
[16]: #Normalize the numeric variables: this is step is important specially when we
      ↪ want to build the logistic regression model
      #as the numeric variables are not normalliy distriputed and the LG required a
      ↪ normalized dataset

      from sklearn.preprocessing import MinMaxScaler
      from sklearn import preprocessing

      for column in dff.columns:
          if dff[column].dtype in ['int64', 'float64']:
              X = dff[column].array.reshape(-1,1)
              #build the scaler model
              scaler = MinMaxScaler()

              # fit using the whole dataset
              scaler.fit(X)

              dff[column]=scaler.transform(dff[column].array.reshape(-1, 1)) # Apply
      ↪ to to the dataset
      display(dff)
```

	age	job	marital	education	default	balance	housing	\
0	0.519481	management	married	tertiary	no	0.092259	yes	
1	0.337662	technician	single	secondary	no	0.073067	yes	
2	0.194805	entrepreneur	married	secondary	no	0.072822	yes	
3	0.376623	blue-collar	married	unknown	no	0.086476	yes	
4	0.194805	unknown	single	unknown	no	0.072812	no	
...	
45206	0.428571	technician	married	tertiary	no	0.080293	no	
45207	0.688312	retired	divorced	primary	no	0.088501	no	
45208	0.701299	retired	married	secondary	no	0.124689	no	
45209	0.506494	blue-collar	married	secondary	no	0.078868	no	
45210	0.246753	entrepreneur	married	secondary	no	0.099777	no	

	loan	contact	day	month	duration	campaign	pdays	previous	\
0	no	unknown	0.133333	may	0.053070	0.000000	0.000000	0.000000	
1	no	unknown	0.133333	may	0.030704	0.000000	0.000000	0.000000	
2	yes	unknown	0.133333	may	0.015453	0.000000	0.000000	0.000000	
3	no	unknown	0.133333	may	0.018707	0.000000	0.000000	0.000000	
4	no	unknown	0.133333	may	0.040260	0.000000	0.000000	0.000000	
...	
45206	no	cellular	0.533333	nov	0.198658	0.032258	0.000000	0.000000	
45207	no	cellular	0.533333	nov	0.092721	0.016129	0.000000	0.000000	
45208	no	cellular	0.533333	nov	0.229158	0.064516	0.212156	0.010909	

```

45209    no    telephone    0.533333    nov    0.103294    0.048387    0.000000    0.000000
45210    no      cellular    0.533333    nov    0.073404    0.016129    0.216743    0.040000

```

```

      poutcome    y
0      unknown    no
1      unknown    no
2      unknown    no
3      unknown    no
4      unknown    no
...          ...   ...
45206    unknown    yes
45207    unknown    yes
45208    success    yes
45209    unknown    no
45210      other    no

```

[45211 rows x 17 columns]

```

[17]: #from sklearn import preprocessing
      #d = preprocessing.normalize(dff.select_dtypes('int64'))
      #names=dff.select_dtypes('int64').columns
      #scaled_df = pd.DataFrame(d, columns=names)
      #scaled_df

```

label encoder for variable: y

```

[18]: le = preprocessing.LabelEncoder()
      le.fit(dff["y"])
      list(le.classes_)
      dff["y"]=le.transform(dff["y"])
      # we converted the dependent variable to be 0,1 where 0 is for "no" and 1 is_
      ↪for "yes" which is the desired output

```

Get dummies

```

[19]: #get dummies for the rest of categorical variables:
      df_cat = dff.select_dtypes('object')
      dff = pd.get_dummies(dff, df_cat.columns, drop_first = False)
      dff.info()
      # we got 52 coulmns (variables)

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 52 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   45211 non-null  float64
1   balance               45211 non-null  float64

```

2	day	45211	non-null	float64
3	duration	45211	non-null	float64
4	campaign	45211	non-null	float64
5	pdays	45211	non-null	float64
6	previous	45211	non-null	float64
7	y	45211	non-null	int64
8	job_admin.	45211	non-null	uint8
9	job_blue-collar	45211	non-null	uint8
10	job_entrepreneur	45211	non-null	uint8
11	job_housemaid	45211	non-null	uint8
12	job_management	45211	non-null	uint8
13	job_retired	45211	non-null	uint8
14	job_self-employed	45211	non-null	uint8
15	job_services	45211	non-null	uint8
16	job_student	45211	non-null	uint8
17	job_technician	45211	non-null	uint8
18	job_unemployed	45211	non-null	uint8
19	job_unknown	45211	non-null	uint8
20	marital_divorced	45211	non-null	uint8
21	marital_married	45211	non-null	uint8
22	marital_single	45211	non-null	uint8
23	education_primary	45211	non-null	uint8
24	education_secondary	45211	non-null	uint8
25	education_tertiary	45211	non-null	uint8
26	education_unknown	45211	non-null	uint8
27	default_no	45211	non-null	uint8
28	default_yes	45211	non-null	uint8
29	housing_no	45211	non-null	uint8
30	housing_yes	45211	non-null	uint8
31	loan_no	45211	non-null	uint8
32	loan_yes	45211	non-null	uint8
33	contact_cellular	45211	non-null	uint8
34	contact_telephone	45211	non-null	uint8
35	contact_unknown	45211	non-null	uint8
36	month_apr	45211	non-null	uint8
37	month_aug	45211	non-null	uint8
38	month_dec	45211	non-null	uint8
39	month_feb	45211	non-null	uint8
40	month_jan	45211	non-null	uint8
41	month_jul	45211	non-null	uint8
42	month_jun	45211	non-null	uint8
43	month_mar	45211	non-null	uint8
44	month_may	45211	non-null	uint8
45	month_nov	45211	non-null	uint8
46	month_oct	45211	non-null	uint8
47	month_sep	45211	non-null	uint8
48	poutcome_failure	45211	non-null	uint8
49	poutcome_other	45211	non-null	uint8

```

50 poutcome_success      45211 non-null  uint8
51 poutcome_unknown      45211 non-null  uint8
dtypes: float64(7), int64(1), uint8(44)
memory usage: 4.7 MB

```

4 Split the Data

```

[20]: #The split technique that have been chosen is holdout split (Test set is 33.
      ↪3% and Training set is 66.7%)
      #as this is the most popular way to use, only needs to be run once so has lower
      ↪computational costs.
      #Many of the researchers on similar topic used this technique and ratio
      #Since this method might lead to an overfitting after oversampling technique or
      ↪underfitting after undersampling technique.
      #The performance will be checked again and average will be taken, after applying
      ↪oversampling/undersampling and cross validation

import numpy as np
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test= train_test_split(dff.
      ↪drop(["y"],axis="columns"),dff.y,shuffle=True, random_state=5,test_size=.333)

```

```

[21]: len(x_train)#number or rows in train set

```

```

[21]: 30155

```

```

[22]: len(x_test)#number or rows in test set

```

```

[22]: 15056

```

Convert the y variable in the train and test set to flattened array(1D array with all the input-array elements and with the same type as it).

```

[23]: y_train = y_train.ravel()
      y_test = y_test.ravel()

```

5 Modles Building

The chosen models are Random Forest, Decision Tree, and Logistic Regression. The reason behind choosing these algorithms is that they are the most popular models for a binary classification problem. As what mentioned in the literature review these models were the most common models to solve a binary classification problem on dataset similar/exact to the dataset of this project.

Random forest is the candidate to be the best model perfomance as it was the best in many reviewed related articles.

In addition to that, many studies emphasized that ensemble models performs better than other models.

5.1 Imbalanced Data

1. In this stage the models have been built on the Imbalanced dataset, in order to compare the performance results before and after balancing the data. later in this script, an application on over sampling (SMOTE) and undersampling techniques (Random Under Sampling) will be applied and compared.
2. The models have been tuned based on some parameters: A. Random Forest tuned parameters are: Max depth and the n_estimators (number of trees) B. Decision Tree tuned parameters are: Max depth
3. We focused on optimizing F1-Score metric as it was the smallest value in the confusion matrix in the first run time. please note that the tuning process also have been applied to optimize the Accuracy and the AUROC but it didnt improve as expected also it led to a decrease in the MCC while improving the F1-score led to an improvement in MCC.

5.1.1 Random Forest Hyperparameter Tuning (based on max depth, number of trees and F1-Score)

```
[ ]: '''Grid search (GridSearchCV) is arguably the most basic hyperparameter tuning
    ↪method. With this technique,
    we simply build a model for each possible combination of all of the
    ↪hyperparameter values provided,
    evaluating each model, and selecting the architecture which produces the best
    ↪results.
```

```
For more information about GridSearchCV: https://scikit-learn.org/stable/
    ↪modules/generated/sklearn.model_selection.GridSearchCV.html'''
```

```
[ ]: 'Grid search (GridSearchCV) is arguably the most basic hyperparameter tuning
method. With this technique, \nwe simply build a model for each possible
combination of all of the hyperparameter values provided, \nevaluating each
model, and selecting the architecture which produces the best results.\n\nFor
more information about GridSearchCV: https://scikit-
learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html'
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(random_state=1)
```

```
[ ]: #This step has been run repeatedly on different ranges of the model parameters
    ↪in order to tune them on the right values.
import numpy as np
from sklearn.model_selection import GridSearchCV
```

```

max_depth_range=np.arange(25,30,1)
n_estimators_range=np.arange(130,180,30)
param_grid=dict(max_depth=max_depth_range,n_estimators=n_estimators_range)
Urfgrid=GridSearchCV(estimator=rf,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)

```

```
[ ]: Urfgrid.fit(x_train,y_train)
```

```
[ ]: GridSearchCV(cv=3, estimator=RandomForestClassifier(random_state=1), n_jobs=-1,
                param_grid={'max_depth': array([25, 26, 27, 28, 29]),
                            'n_estimators': array([130, 160])},
                scoring='f1')
```

```
[ ]: print("the best parameters are %s with a score of %0.2f" % (Urfgrid.
    ↳best_params_,Urfgrid.best_score_))#Mean cross-validated score of the
    ↳best_estimator
#the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params

```

the best parameters are {'max_depth': 29, 'n_estimators': 130} with a score of 0.47

```
[ ]: #data frame of grid search parameters and their F1 Scores scores
import pandas as pd

grid_results = pd.concat([pd.DataFrame(Urfgrid.cv_results_["params"]),
                          pd.DataFrame(Urfgrid.cv_results_["mean_test_score"],
    ↳columns=["f1"])]),axis=1)
grid_results.head()

#preparing data for making contour plots
grid_contour = grid_results.groupby(['max_depth','n_estimators']).mean()
grid_contour

#pivot data:
grid_reset=grid_contour.reset_index()
grid_reset.columns=["max_depth","n_estimator","f1"]
grid_pivot=grid_reset.pivot("max_depth","n_estimator")

#assigning the pivoted data into the respective x,y, and z variables
x=grid_pivot.columns.levels[1].values
y=grid_pivot.index.values
z=grid_pivot.values

#2D contour plot
import plotly.graph_objects as go

# X and Y axes labels

```

```

layout = go.Layout(
    xaxis=go.layout.XAxis(
        title=go.layout.xaxis.Title(
            text='n_estimators')
    ),
    yaxis=go.layout.YAxis(
        title=go.layout.yaxis.Title(
            text='max_depth')
    )

fig = go.Figure(data = [go.Contour(z=z, x=x, y=y)], layout=layout )

fig.update_layout(title='Imbalanced- Random Forest Hyperparameter tuning',
    ↪autosize=False,
    width=500, height=500,
    margin=dict(l=65, r=50, b=65, t=90))

fig.show()

```

5.1.2 Decision Tree Hyperparameter Tuning (based on max depth and F1-Score)

```

[ ]: from sklearn.tree import DecisionTreeClassifier
    dt=DecisionTreeClassifier(random_state=1)

[ ]: #This step has been run repeatedly on different ranges of the model parameters
    ↪in order to tune them on the right values.
    import numpy as np
    from sklearn.model_selection import GridSearchCV

    max_depth_range=np.arange(1,10,1)
    param_grid=dict(max_depth=max_depth_range)
    Udtgrid=GridSearchCV(estimator=dt,param_grid=param_grid,scoring="f1",cv=10,n_jobs=-1)

[ ]: Udtgrid.fit(x_train,y_train)

[ ]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(random_state=1), n_jobs=-1,
    param_grid={'max_depth': array([1, 2, 3, 4, 5, 6, 7, 8, 9])},
    scoring='f1')

[ ]: print("the best parameters are %s with a score of %0.2f" % (Udtgrid.
    ↪best_params_,Udtgrid.best_score_))#Mean cross-validated score of the
    ↪best_estimator
    #the score is the average of all cv folds for a single combination of the
    ↪parameters you specify in the tuned_params

```

the best parameters are {'max_depth': 9} with a score of 0.48

5.1.3 Logistic Regression Hyperparameter Tuning (based on Max Iteration, Solver and F1-Score)

```
[ ]: from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(random_state=1,n_jobs=-1)

[ ]: #This step has been run repeatedly on different ranges of the model parameters
→in order to tune them on the right values.
import numpy as np
from sklearn.model_selection import GridSearchCV

max_iter=np.arange(100,1000,100)

param_grid=dict(max_iter=max_iter,solver=['liblinear', 'lbfgs'])
Ulgrrgrid=GridSearchCV(estimator=lgr,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)
# -1 means using all processors. This means that the n_jobs parameter can be
→used to distribute and exploit all the CPUs available in the local computer

[ ]: Ulgrrgrid.fit(x_train,y_train)

[ ]: GridSearchCV(cv=3, estimator=LogisticRegression(n_jobs=-1, random_state=1),
                n_jobs=-1,
                param_grid={'max_iter': array([100, 200, 300, 400, 500, 600, 700,
800, 900]),
                            'solver': ['liblinear', 'lbfgs']}},
                scoring='f1')

[ ]: print("the best parameters are %s with a score of %0.2f" % (Ulgrrgrid.
    →best_params_,Ulgrrgrid.best_score_))#Mean cross-validated score of the
    →best_estimator
    #the score is the average of all cv folds for a single combination of the
    →parameters you specify in the tuned_params
```

the best parameters are {'max_iter': 100, 'solver': 'lbfgs'} with a score of 0.42

5.1.4 1. Imbalanced-Train the model and calculate the run time of training the model

```
[24]: #From the above we can conclude that the best parameters are:

#Random Forest:
U_RF_max_depth= 29 #Is the max depth for each tree in the random forest
U_RF_n_estimators= 130 #Is the number of trees in the random forest

# Decision Tree:
U_DT_max_depth=9 #Is the max depth for the tree
```

```

#Logistic Regression:
U_lgr_max_iter=100 # max_iter is the maximum number of iterations taken for the
↳ solvers to converge.

U_lgr_solver="lbfgs" #Solver: is the algorithm to use in the optimization
↳ problem

#These values may differ slightly on each run for tuning.
#Please note that I have saved the parameters values (after tuning) of each
↳ model in the above variables in order to make the run time shorter everytime
↳ we need to train the model.

```

```

[25]: #Random Forest Training
np.random.seed(1)
import time
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=U_RF_n_estimators, max_depth=
↳ U_RF_max_depth,random_state=1)

startrf=time.time()
rf.fit(x_train,y_train)
endrf=time.time()

#Decision Tree Training
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(max_depth=U_DT_max_depth,random_state=1)

startdt=time.time()
dt.fit(x_train,y_train)
enddt=time.time()

#Linear Regression Training
from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(max_iter=U_lgr_max_iter, solver=U_lgr_solver,n_jobs=-1)
# -1 means using all processors. This means that the n_jobs parameter can be
↳ used to distribute and exploit all the CPUs available in the local computer
startlgr=time.time()
lgr.fit(x_train,y_train)
endlgr=time.time()

print("RF Training Run Time [s]:",endrf-startrf)
print("DT Training Run Time [s]:",enddt-startdt)
print("LGR Training Run Time [s]:",endlgr-startlgr)

# The results shows the the faster model is the Decision tree, Logistic
↳ regression, and the Random forest respectively.

```

```
# Random Forest model was the slowest one( which is expected as it is an  
→ensemble model)
```

RF Training Run Time [s]: 5.8421471118927
DT Training Run Time [s]: 0.19419074058532715
LGR Training Run Time [s]: 2.004239559173584

5.1.5 2. Imbalanced- Run Time for prediction

```
[26]: # The run time to predict the test set in each model is:  
np.random.seed(2)  
#Random Forest  
startrf=time.time()  
y_predictedrf=rf.predict(x_test)  
endrf=time.time()  
  
#Decision Tree  
startdt=time.time()  
y_predictdt=dt.predict(x_test)  
enddt=time.time()  
  
#Logistic Regression  
startlgr=time.time()  
y_predictedlgr=lgr.predict(x_test)  
endlgr=time.time()  
  
print("RF Run Time to Predict the testset [s]:",endrf-startrf)  
print("DT Run Time to Predict the testset [s]:",enddt-startdt)  
print("LGR Run Time to Predict the testset [s]:",endlgr-startlgr)  
  
# The results shows the the faster model is the Logistic regression,Decision  
→tree, and the Random forest respectively.  
# Random Forest model was the slowest one( which is expected as it is an  
→ensemble model)
```

RF Run Time to Predict the testset [s]: 0.5863540172576904
DT Run Time to Predict the testset [s]: 0.0069544315338134766
LGR Run Time to Predict the testset [s]: 0.005115032196044922

5.1.6 3. Imbalanced- Confusion Matrix

```
[27]: from matplotlib import pyplot as plt  
from sklearn.metrics import recall_score, make_scorer, confusion_matrix,  
→classification_report,ConfusionMatrixDisplay  
import matplotlib
```

```

import matplotlib.pyplot as plt
import seaborn as sns
class Colors:
    Gray = "#5d5d5d"
    LightGray = "#fafafa"
    Black = "#000000"
    White = "#FFFFFF"
    Teal = "#008080"
    Aquamarine = "#76c8c8"
    Blue = "#2596be"
    LightCyan = "#badbdb"
    WhiteSmoke = "#dedad2"
    Cream = "#e4bcad"
    PeachPuff = "#df979e"
    HotPink = "#d7658b"
    DeepPink = "#c80064"
    LightSeaGreen = "#20B2AA"
    DarkGray = "#464144"

#Confusion Matrix of Random Forest
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmrf = confusion_matrix(y_test,y_predictedrf)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmrf, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↪Subscribed'],annot_kws={"fontsize":14})
print("Random Forest \n",classification_report(y_test, y_predictedrf))
ax0.text(0,-0.75,'Random Forest_
↪Performance',fontsize=18,fontweight='bold',fontfamily='serif')

#Confusion Matrix of Decision Tree
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmdt = confusion_matrix(y_test,y_predicteddt)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]

```

```

colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmdt, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↳Subscribed'],annot_kws={"fontsize":14})
print("Decision Tree\n",classification_report(y_test, y_predicteddt))
ax0.text(0,-0.75,'Decision Tree_
↳Performance',fontsize=18,fontweight='bold',fontfamily='serif')

#Confusion Matrix of Logistic Regression
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmlgr = confusion_matrix(y_test,y_predictedlgr)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↳Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmlgr, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↳Subscribed'],annot_kws={"fontsize":14})
print("Logistic Regression\n",classification_report(y_test, y_predictedlgr))
ax0.text(0,-0.75,'Logistic Regression_
↳Performance',fontsize=18,fontweight='bold',fontfamily='serif')

```

Random Forest



	precision	recall	f1-score	support
0	0.92	0.97	0.95	13276
1	0.67	0.39	0.50	1780
accuracy			0.91	15056
macro avg	0.80	0.68	0.72	15056
weighted avg	0.89	0.91	0.89	15056

Decision Tree

	precision	recall	f1-score	support
0	0.92	0.97	0.94	13276
1	0.62	0.39	0.48	1780
accuracy			0.90	15056
macro avg	0.77	0.68	0.71	15056
weighted avg	0.89	0.90	0.89	15056

Logistic Regression					
	precision	recall	f1-score	support	
0	0.92	0.98	0.95	13276	
1	0.67	0.33	0.44	1780	
accuracy			0.90	15056	
macro avg	0.79	0.66	0.70	15056	
weighted avg	0.89	0.90	0.89	15056	

[27]: Text(0, -0.75, 'Logistic Regression Performance')



Logistic Regression Performance		
Actual Non-Subscribed	12986	290
Actual Subscribed	1188	592
	Predicted Non-Subscribed	Predicted Subscribed

```
[28]: #checking
tn, fp, fn, tp = cmrf.ravel()
tn, fp, fn, tp
```

```
[28]: (12927, 349, 1079, 701)
```

5.1.7 4. Imbalanced- Calculate the AUROC

1. Higher is better
2. Between 0 and 1

```
[29]: #to get the probabilities of the prediction for each instance in the test set
np.random.seed(1)
rf_probs= rf.predict_proba(x_test)#random forest
dt_probs= dt.predict_proba(x_test)#decision tree
lgr_probs= lgr.predict_proba(x_test)#logistic regression

#to keep the probabilities of positive outcomes

#random forest
rf_probs=rf_probs[:,1]
rf_probs

#decision tree
dt_probs=dt_probs[:,1]
dt_probs

#logistic regression
lgr_probs=lgr_probs[:,1]
lgr_probs
```

```

#Print AUROC score
from sklearn.metrics import roc_curve, roc_auc_score
rf_auc=roc_auc_score(y_test,rf_probs)#random forest
dt_auc=roc_auc_score(y_test,dt_probs)#decision tree
lgr_auc=roc_auc_score(y_test,lgr_probs)#logistic regression

print("Random Forest: AUROC= %.3f" % (rf_auc))#random forest
print("Decision Tree: AUROC= %.3f" % (dt_auc))#decision tree
print("Logistic Regression: AUROC= %.3f" % (lgr_auc))#logistic regression

#Calculate and plot ROC curve
'''roc_curve() returns three arrays, first one is false positive rate(fpr),
→second one is true positive rate(tpr) and the third one is the threshold (_)
#for more information: https://scikit-learn.org/stable/modules/generated/
→sklearn.metrics.roc_curve.html'''
rf_fpr,rf_tpr,_=roc_curve(y_test,rf_probs)#random forest
dt_fpr,dt_tpr,_=roc_curve(y_test,dt_probs)#decision tree
lgr_fpr,lgr_tpr,_=roc_curve(y_test,lgr_probs)#logistic regression

#plot each model:
import matplotlib.pyplot as plt
plt.plot(rf_fpr,rf_tpr,marker=".",label="Imbalanced- Random Forest (AUROC= %0.
→3f)" % rf_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.plot(dt_fpr,dt_tpr,marker=".",label="Imbalanced- Decision Tree (AUROC= %0.
→3f)" % dt_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

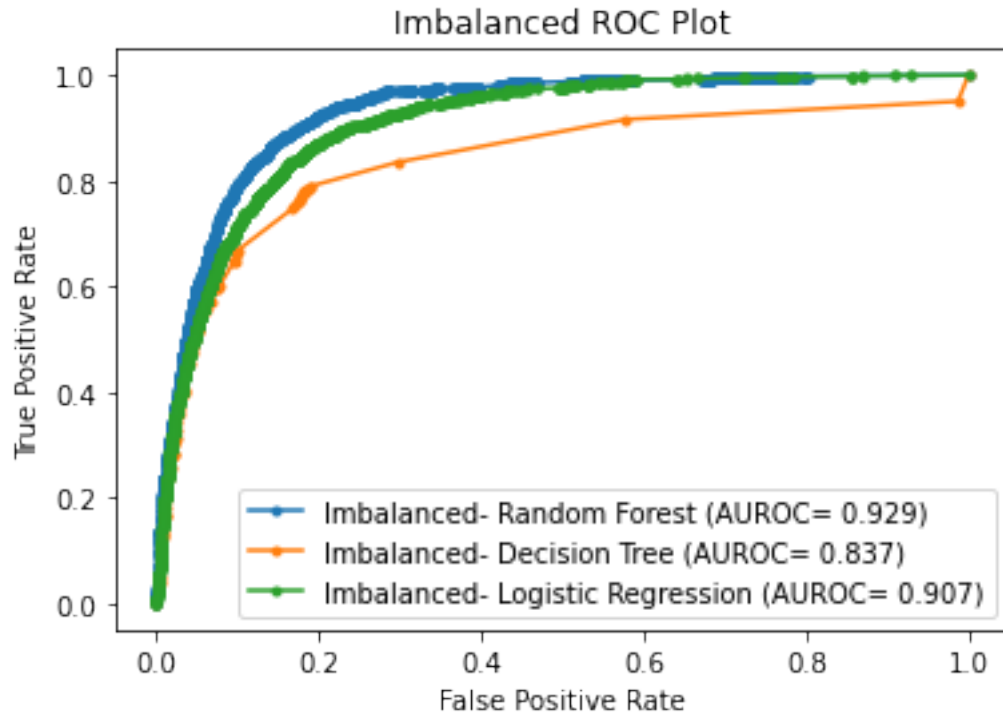
plt.plot(lgr_fpr,lgr_tpr,marker=".",label="Imbalanced- Logistic Regression,
→(AUROC= %0.3f)" % lgr_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.title("Imbalanced ROC Plot")
plt.legend()
plt.show()

#The results show the best AUROC if for Random forest with 0.93 score followed,
→by Logisitic regression and Decision tree respectively.

```

Random Forest: AUROC= 0.929
 Decision Tree: AUROC= 0.837
 Logistic Regression: AUROC= 0.907



5.1.8 5. Imbalanced- Summary (Accuracy, Recall, specificity, AUROC, MCC, Brier_score_loss, Run_time) for the models

```
[30]: np.random.seed(1)
from sklearn.metrics import roc_curve,
    ↳roc_auc_score,accuracy_score,recall_score,matthews_corrcoef,brier_score_loss,f1_score
from sklearn import tree, linear_model,ensemble
MLA = [tree.DecisionTreeClassifier(max_depth=U_DT_max_depth,random_state=1),
        linear_model.
    ↳LogisticRegression(max_iter=U_lgr_max_iter,solver=U_lgr_solver,n_jobs=-1),
        ensemble.RandomForestClassifier(n_estimators=U_RF_n_estimators,
    ↳max_depth=U_RF_max_depth,random_state=1)]
import time
name = []
Accuracy = []
Specificity=[]
Sensetivity=[]
alg_auc=[]
Mcc=[]
br_s=[]
f1score=[]
T_time=[]
```

```

P_time=[]
for alg in MLA:
    name.append(alg.__class__.__name__)
    start=time.time()
    alg.fit(x_train,y_train)
    end=time.time()
    startp=time.time()
    y_predicted=alg.predict(x_test)
    endp=time.time()
    Accuracy.append(accuracy_score(y_test,y_predicted))
    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))
    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))
    Mcc.append(matthews_corrcoef(y_test, y_predicted))
    br_s.append(brier_score_loss(y_test, y_predicted))
    f1score.append(f1_score(y_test,y_predicted))
    alg_probs= alg.predict_proba(x_test)
    alg_probs=alg_probs[:,1]
    alg_auc.append(roc_auc_score(y_test,alg_probs))
    T_time.append(end-start)
    P_time.append(endp-startp)

comparison = pd.DataFrame({"name": name, "Accuracy": Accuracy,"Sensetivity":
    ↳Sensetivity,"Specificity":Specificity,"AUC":alg_auc,"MCC":Mcc,"Brier_score":
    ↳br_s,
                           "F1Score":f1score,"TrainingTime_s":
    ↳T_time,"Prediction Time [s]": P_time})
display(comparison)

#From the below results we can conclude that the random forest is slightly
    ↳better than other models especially
#If we compared them by the Brior_score, Matthew's correlation coefficient
    ↳(MCC), Accuracy, and Specificity.
#MCC: Recently alot of researches prove that it is a better measurement
    ↳especially for binary classification problems along with the brier score.

```

	name	Accuracy	Sensetivity	Specificity	AUC \
0	DecisionTreeClassifier	0.899243	0.394382	0.966933	0.837322
1	LogisticRegression	0.901833	0.332584	0.978156	0.907439
2	RandomForestClassifier	0.905154	0.393820	0.973712	0.929200

	MCC	Brier_score	F1Score	TrainingTime_s	Prediction Time [s]
0	0.440803	0.100757	0.480657	0.223819	0.006711
1	0.427229	0.098167	0.444778	1.102673	0.009281
2	0.465894	0.094846	0.495406	5.836085	0.515099

5.1.9 6. Imbalanced- Check Over/Underfitting

Brier Score was used to evaluate the model fitting

The reason behind using the Brier score metric is that Brier score measures the mean squared difference between the predicted probability and the actual outcome. The Brier score always takes on a value between zero and one, since this is the largest possible difference between a predicted probability (which must be between zero and one) and the actual outcome (which can take on values of only 0 and 1). It can be decomposed is the sum of refinement loss and calibration loss. The Brier score is appropriate for binary and categorical outcomes that can be structured as true or false, but is inappropriate for ordinal variables which can take on three or more values

Brier score is actually 1-Accuracy

”The Brier score is a strictly proper scoring function that is equivalent to the mean squared error” (Chicco, D., Warrens, M. J., & Jurman, G. (2021)).

*The smaller the Brier score loss, the better, hence the naming with “loss”.

1. If the performance of the model on the training set is poor that means the model is underfitting to the test data.
2. If the performance of the model on the test set is poor that means the model is overfitting to the training data.

```
[ ]: #repeated kfold validation
from numpy import mean
from numpy import std
from sklearn.metrics import roc_curve, \
    roc_auc_score, accuracy_score, recall_score, brier_score_loss
from sklearn import tree, linear_model, ensemble
from sklearn.model_selection import validation_curve
from sklearn.model_selection import KFold
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score
import numpy as np
np.random.seed(0)
import matplotlib.pyplot as plt
plt.style.use("ggplot")
```

```
[ ]: #Decision Tree
# prepare the cross-validation procedure
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
U_DT_list_training_BrierScore=[]
U_DT_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
```

```

model=tree.DecisionTreeClassifier(max_depth=U_DT_max_depth,random_state=1)
model.fit(X_train,Y_train)

y_train_data_pred=model.predict(X_train)
y_test_data_pred=model.predict(X_test)

U_DT_fold_training_BrierScore=brier_score_loss(Y_train,y_train_data_pred)
U_DT_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

U_DT_list_training_BrierScore.append(U_DT_fold_training_BrierScore)
U_DT_list_testing_BrierScore.append(U_DT_fold_testing_BrierScore)

```

```

[ ]: print("DT_Test_BScore_Mean:", mean(U_DT_list_testing_BrierScore))
      print("DT_Train_BScore_Mean:", mean(U_DT_list_training_BrierScore))

```

```

DT_Test_BScore_Mean: 0.09962907548611827
DT_Train_BScore_Mean: 0.07913020226274736

```

```

[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing_
      ↪phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_DT_list_training_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Decision Tree \n Training BrierScore across folds")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_DT_list_testing_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title("Decision Tree \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

```

```

#From the below graphs if we looked at the the first graph "Decision Tree_
↪-Training BrierScore across folds" we can see the the models fits really_
↪well as the BrierScore
#is around 0.08.
#in the second graph "Decision Tree -Testing BrierScore across folds"it also_
↪fits well as the BrierScore on the test set is around 0.1.
#That means also that we have a fitted model and stable one.

```



```
[ ]: #Logistic Regression
# prepare the cross-validation procedure
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
U_lgr_list_training_BrierScore=[]
U_lgr_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index:],X.iloc[test_index:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
    model=linear_model.LogisticRegression(max_iter=U_lgr_max_iter,solver=U
    ↪U_lgr_solver,random_state=1,n_jobs=-1)
    model.fit(X_train,Y_train)

    y_train_data_pred=(model.predict_proba(X_train))[:,1]
    y_test_data_pred= (model.predict_proba(X_test))[:,1]

    U_lgr_fold_training_BrierScore=brier_score_loss(Y_train,y_train_data_pred)
    U_lgr_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

    U_lgr_list_training_BrierScore.append(U_lgr_fold_training_BrierScore)
    U_lgr_list_testing_BrierScore.append(U_lgr_fold_testing_BrierScore)
```

```
[ ]: print("LG_Test_BScore_Mean:", mean(U_lgr_list_testing_BrierScore))
print("LG_Train_BScore_Mean:", mean(U_lgr_list_training_BrierScore))
```

```
LG_Test_BScore_Mean: 0.07132303555637362
LG_Train_BScore_Mean: 0.07100383987335479
```


```
[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
    ↪phase
plt.figure(figsize = (10, 4))
```

```

plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_lgr_list_training_BrierScore).
        ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Logistic BrierScore \n Training BrierScore across folds")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_lgr_list_testing_BrierScore).
        ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title("Logistic Regression \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

```

*#From the below graphs if we looked at the the first graph "Logistic Regression ↪
 ↪-Training BrierScore across folds" we can see the the models fits really ↪
 ↪well
 #as the BrierScore is stable and very small (about 0.07).
 #In the second graph "Logistic Regression -Testing BrierScore across folds"it ↪
 ↪also fits well as the BrierScore on the test set is also around 0.07.
 #That means also that we have a fitted model and stable one.*



```

[ ]: #Random Forest
cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)
U_RF_list_training_BrierScore=[]
U_RF_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(dff):

```

```

X_train,X_test=X.iloc[train_index:],X.iloc[test_index:]
Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
model=ensemble.RandomForestClassifier(n_estimators=U_RF_n_estimators,
↳max_depth=U_RF_max_depth,random_state=1)
model.fit(X_train,Y_train)

y_train_data_pred=model.predict(X_train)
y_test_data_pred=model.predict(X_test)

U_RF_fold_training_BrierScore=brier_score_loss(Y_train,y_train_data_pred)
U_RF_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

U_RF_list_training_BrierScore.append(U_RF_fold_training_BrierScore)
U_RF_list_testing_BrierScore.append(U_RF_fold_testing_BrierScore)

```

```

[ ]: print("RF_Test_BScore_Mean:", mean(U_RF_list_testing_BrierScore))
print("RF_Train_BScore_Mean:", mean(U_RF_list_training_BrierScore))

```

```

RF_Test_BScore_Mean: 0.09349489758068094
RF_Train_BScore_Mean: 0.0005267483980666975

```

```


[ ]: #plot the BrierScore of the training phase and the accuBrierScoreracy of the
↳testing phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_RF_list_training_BrierScore).
↳ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Random Forest \n Training BrierScore across folds")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(U_RF_list_testing_BrierScore).
↳ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title("Random Forest \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

```

```

#From the below graphs if we looked at the the first graph "Random
↳Forest-Training BrierScore across folds" we can see the the models fits
↳really well as the BrierScore is almost 0
#in the second graph "Random Forest -Testing BrierScore across folds"it also
↳fits well as the BrierScore on the test set is around 0.1.
#That means also that we have a fitted model and stable one.

```



5.2 SMOTE

1. In this stage the models have been built on balanced dataset using SMOTE (Oversampling technique).
2. Sampling strategy parameter in SMOTE function has been tuned for each model in order to optimize the F1 Score.
3. The models have been tuned based on some parameters: A. Random Forest tuned parameters are: Max depth and the n_estimators (number of trees) . B. Decision Tree tuned parameters are: Max depth. C. Logistic Regression tuned parameters are : Max iteration and Solver.
4. We focused on optimizing F1-Score metric as it was the smallest value in the confusion matrix in the first run

5.2.1 1. SMOTE- Hyperparameter Tuning (based on Sampling strategy and F1-Score)

Please note that the sampling strategy is the desired ratio of the number of samples in the minority class over the number of samples in the majority class after resampling.

```
[ ]: import imblearn
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from imblearn.over_sampling import SMOTE

[ ]: #Tuning SMOTE for Random Forest
#This step has been run repeatedly on different ranges of sampling_strategy_
    ↳parameter in order to tune it on the right value.
rfweights=np.linspace(0.55,1,20)
pipe=imblearn.pipeline.make_pipeline(SMOTE(),RandomForestClassifier())
```




```
gsc=GridSearchCV(estimator=pipe,param_grid={"smote__sampling_strategy":
↳rfweights},scoring="f1",cv=3,n_jobs=-1)
Srf_grid_result=gsc.fit(x_train,y_train)
```

```
[ ]: #Random Forest
print("RF-the best parameters are %s with a score of %0.2f" % (Srf_grid_result.
↳best_params_,Srf_grid_result.best_score_))#Mean cross-validated score of the
↳best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":Srf_grid_result.
↳cv_results_["mean_test_score"],"weight":rfweights})
#the score is the average of all cv folds for a single combination of the
↳parameters you specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Random Forest")
```

RF-the best parameters are {'smote__sampling_strategy': 0.8342105263157895} with a score of 0.57

```
[ ]: <AxesSubplot: title={'center': 'Random Forest'}, xlabel='weight'>
```



```
[ ]: #Tuning SMOTE for Decision Tree
#This step has been run repeatedly on different ranges of sampling_strategy
↳parameter in order to tune it on the right value.
dtweights=np.linspace(0.4,1,20)
```

```

pipe=imblearn.pipeline.make_pipeline(SMOTE(),DecisionTreeClassifier())
gsc=GridSearchCV(estimator=pipe,param_grid={"smote__sampling_strategy":
    ↳dtweights},scoring="f1",cv=3,n_jobs=-1)
Sdt_grid_result=gsc.fit(x_train,y_train)

```

```

[ ]: #Decision Tree
print("DT-the best parameters are %s with a score of %0.2f" % (Sdt_grid_result.
    ↳best_params_,Sdt_grid_result.best_score_))#Mean cross-validated score of the
    ↳best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":Sdt_grid_result.
    ↳cv_results_["mean_test_score"],"weight":dtweights})
#the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Decision Tree")

```

DT-the best parameters are {'smote__sampling_strategy': 0.8105263157894738} with a score of 0.49

```

[ ]: <AxesSubplot: title={'center': 'Decision Tree'}, xlabel='weight'>

```



```

[ ]: #Tuning SMOTE for Logistic Regression
#This step has been run repeatedly on different ranges of sampling_strategy
    ↳parameter in order to tune it on the right value.


```

```
lgrweights=np.linspace(0.4,1,20)
pipe=imblearn.pipeline.make_pipeline(SMOTE(),LogisticRegression(n_jobs=-1))
gsc=GridSearchCV(estimator=pipe,param_grid={"smote__sampling_strategy":
    ↳lgrweights},scoring="f1",cv=3,n_jobs=-1)
Slgr_grid_result=gsc.fit(x_train,y_train)
```

```
[ ]: #Logistic Regression
print("LGR-the best parameters are %s with a score of %0.2f" %
    ↳(Slgr_grid_result.best_params_,Slgr_grid_result.best_score_))#Mean
    ↳cross-validated score of the best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":Slgr_grid_result.
    ↳cv_results_["mean_test_score"],"weight":lgrweights})
#the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Logistic Regression")
```

LGR-the best parameters are {'smote__sampling_strategy': 1.0} with a score of 0.56

```
[ ]: <AxesSubplot: title={'center': 'Logistic Regression'}, xlabel='weight'>
```



5.2.2 2. Data Balancing: (SMOTE)

```
[31]: #from the above we can conclude that the best parameters for smote are:

#Random Forest:
RF_SMT_sampling_strategy=0.8342105263157895

# Decision Tree:
DT_SMT_sampling_strategy=0.8105263157894738

# Logistic Regression:
lgr_SMT_sampling_strategy=1

#These values may differ slightly on each run of the previous tuning.
#Please note that I have saved the smote-sampling strategy values (after
→tuning) for each model in the above variables
#in order to make the run time shorter everytime we need to balance the data.
```

```
[32]: #The portions for minority and majority before smote
import numpy as np
unique,count=np.unique(y_train,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[32]: {0: 26646, 1: 3509}
```

```
[33]: #SMOTE resampling-Decision Tree
from imblearn.over_sampling import SMOTE
smdt=SMOTE(sampling_strategy=DT_SMT_sampling_strategy,random_state=101)
dtx_train_res,dyt_train_res=smdt.fit_resample(x_train,y_train)
unique,count=np.unique(dyt_train_res,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[33]: {0: 26646, 1: 21597}
```

```
[34]: #SMOTE resampling-Random Forest
smrf=SMOTE(sampling_strategy=RF_SMT_sampling_strategy,random_state=12)
rfx_train_res,rfy_train_res=smrf.fit_resample(x_train,y_train)
unique,count=np.unique(rfy_train_res,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[34]: {0: 26646, 1: 22228}
```

```
[35]: #SMOTE resampling-Logistic Regression
smtlgr=SMOTE(sampling_strategy=lgr_SMT_sampling_strategy,random_state=13)
```

```
lgrx_train_res,lgry_train_res=smtlgr.fit_resample(x_train,y_train)
unique,count=np.unique(lgry_train_res,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[35]: {0: 26646, 1: 26646}
```

5.2.3 3. SMOTE- Random Forest Hyperparameter Tuning (based on max depth, number of trees and F1-Score)

```
[ ]: from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(random_state=2)

[ ]: import numpy as np
from sklearn.model_selection import GridSearchCV
#This step has been run repeatedly on different ranges of the model parameters
→in order to tune them on the right values.
max_depth_range=np.arange(10,35,1)
n_estimators_range=np.arange(120,200,20)
param_grid=dict(max_depth=max_depth_range,n_estimators=n_estimators_range)
rfgridSmote=GridSearchCV(estimator=rf,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)

[ ]: rfgridSmote.fit(rfx_train_res,rfy_train_res)

[ ]: print("the best parameters are %s with a score of %0.2f" % (rfgridSmote.
    →best_params_,rfgridSmote.best_score_))#Mean cross-validated score of the
    →best_estimator
#the score is the average of all cv folds for a single combination of the
    →parameters you specify in the tuned_params
```

the best parameters are {'max_depth': 32, 'n_estimators': 120} with a score of 0.93

```
[ ]: #data frame of grid search parameters and their accuracy scores
import pandas as pd

grid_results = pd.concat([pd.DataFrame(rfgridSmote.cv_results_["params"]),
                             pd.DataFrame(rfgridSmote.
    →cv_results_["mean_test_score"], columns=["f1"])],axis=1)
grid_results.head()

#preparing data for making contour plots
grid_contour = grid_results.groupby(['max_depth','n_estimators']).mean()
grid_contour
```

```

#pivot data:
grid_reset=grid_contour.reset_index()
grid_reset.columns=["max_depth","n_estimator","f1"]
grid_pivot=grid_reset.pivot("max_depth","n_estimator")

#assigning the pivoted data into the respective x,y, and z variables
x=grid_pivot.columns.levels[1].values
y=grid_pivot.index.values
z=grid_pivot.values

#2D contour plot
import plotly.graph_objects as go

# X and Y axes labels
layout = go.Layout(
    xaxis=go.layout.XAxis(
        title=go.layout.xaxis.Title(
            text='n_estimators')
        ),
    yaxis=go.layout.YAxis(
        title=go.layout.yaxis.Title(
            text='max_depth')
    ) )

fig = go.Figure(data = [go.Contour(z=z, x=x, y=y)], layout=layout )

fig.update_layout(title='Hyperparameter tuning', autosize=False,
                    width=500, height=500,
                    margin=dict(l=65, r=50, b=65, t=90))

fig.show()

```

5.2.4 4. SMOTE- Decision Tree Hyperparameter Tuning (based on max depth and F1-Score)

```

[ ]: from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(random_state=2)

```

```

[ ]: import numpy as np
from sklearn.model_selection import GridSearchCV
#This step has been run repeatedly on different ranges of the model parameters_
↳ in order to tune them on the right values.
max_depth_range=np.arange(1,20,1)
param_grid=dict(max_depth=max_depth_range)
dtgridSmote=GridSearchCV(estimator=dt,param_grid=param_grid,scoring="f1",cv=10,n_jobs=-1)

```

```
[ ]: dtgridSmote.fit(dtx_train_res, dty_train_res)
```

```
[ ]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(random_state=2), n_jobs=-1,
                  param_grid={'max_depth': array([ 1, 2, 3, 4, 5, 6, 7, 8, 9,
10, 11, 12, 13, 14, 15, 16, 17,
18, 19])},
                  scoring='f1')
```

```
[ ]: print("the best parameters are %s with a score of %0.2f" % (dtgridSmote.
    ↳ best_params_, dtgridSmote.best_score_)) #Mean cross-validated score of the
    ↳ best_estimator
#the score is the average of all cv folds for a single combination of the
    ↳ parameters you specify in the tuned_params
```

the best parameters are {'max_depth': 17} with a score of 0.89

5.2.5 3. SMOTE- Logistic Regression Hyperparameter Tuning (based on Max Iteration, Solver and F1-Score)

```
[ ]: from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(random_state=10)
```

```
[ ]: #This step has been run repeatedly on different ranges of the model parameters
    ↳ in order to tune them on the right values.
import numpy as np
from sklearn.model_selection import GridSearchCV

max_iter=np.arange(1000,2100,500)
param_grid=dict(max_iter=max_iter,solver=['liblinear', 'lbfgs'])

lgrgridSmote=GridSearchCV(estimator=lgr,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)
```

```
[ ]: lgrgridSmote.fit(lgrx_train_res, lgry_train_res)
```

```
[ ]: GridSearchCV(cv=3, estimator=LogisticRegression(random_state=10), n_jobs=-1,
                  param_grid={'max_iter': array([1000, 1500, 2000]),
                              'solver': ['liblinear', 'lbfgs']},
                  scoring='f1')
```

```
[ ]: print("the best parameters are %s with a score of %0.2f" % (lgrgridSmote.
    ↳ best_params_, lgrgridSmote.best_score_)) #Mean cross-validated score of the
    ↳ best_estimator
#the score is the average of all cv folds for a single combination of the
    ↳ parameters you specify in the tuned_params.
#The better Solver after resampling the data using SMOTE is "liblinear".
#after oversampling the data "liblinear" becomes better choice.
```

the best parameters are {'max_iter': 1000, 'solver': 'liblinear'} with a score of 0.89

5.2.6 5. SMOTE- Train the model and calculate the run time of training the model

[36]: *#from the above we can conclude that the best parameters are:*

```
#Random Forest:
SMT_RF_max_depth= 32
SMT_RF_n_estimators= 120

# Decision Tree:
SMT_DT_max_depth=17

#logistic Regression:
SMT_lgr_max_iter=1000
SMT_lgr_solver="liblinear"

#These values may differ slightly on each run of the previous tunning.
#Please note that I have saved the parameters values (after tuning) of each
→model in the above variables in order
#to make the run time shorter everytime we need to train the model.'
```

[37]: *#RF*

```
import time
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=SMT_RF_n_estimators, max_depth=
→SMT_RF_max_depth,random_state=2)

startrf=time.time()
rf.fit(rfx_train_res,rfy_train_res)
endrf=time.time()

#DT
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(max_depth=SMT_DT_max_depth,random_state=2)

startdt=time.time()
dt.fit(dtx_train_res,dy_train_res)
enddt=time.time()

#LGR
from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(max_iter=SMT_lgr_max_iter,solver=SMT_lgr_solver,n_jobs=-1,random_state=
startlgr=time.time()
```



```
lgr.fit(lgrx_train_res,lgry_train_res)
endlgr=time.time()

print("RF Training Run Time [s]:",endrf-startrf)
print("DT Training Run Time [s]:",enddt-startdt)
print("LGR Training Run Time [s]:",endlgr-startlgr)
# The results shows the the faster model is the Logistic Regression, Decision
→tree, and Random Forest respectively.
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:1153:
UserWarning: 'n_jobs' > 1 does not have any effect when 'solver' is set to
'liblinear'. Got 'n_jobs' = 2.
  warnings.warn(

RF Training Run Time [s]: 10.67328691482544
DT Training Run Time [s]: 0.616220235824585
LGR Training Run Time [s]: 0.7037904262542725
```

5.2.7 6. SMOTE- Run Time for prediction

```
[38]: # the run time to predict the test set
startrf=time.time()
y_predictedrf=rf.predict(x_test)
endrf=time.time()

startdt=time.time()
y_predicteddt=dt.predict(x_test)
enddt=time.time()

startlgr=time.time()
y_predictedlgr=lgr.predict(x_test)
endlgr=time.time()

print("RF Run Time to Predict the testset [s]:",endrf-startrf)
print("DT Run Time to Predict the testset [s]:",enddt-startdt)
print("LGR Run Time to Predict the testset [s]:",endlgr-startlgr)

# The results shows the the faster model is the Logistic regression, Decision
→tree, and the Random forest respectively.
# Random Forest model was the slowest one( which is expected as it is an
→ensemble model)
```

```
RF Run Time to Predict the testset [s]: 0.5524380207061768
DT Run Time to Predict the testset [s]: 0.008587360382080078
LGR Run Time to Predict the testset [s]: 0.0053730010986328125
```

5.2.8 7. SMOTE- Confusion Matrix

```
[39]: from sklearn.metrics import recall_score, make_scorer, confusion_matrix, \
      ↪ classification_report, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns

class Colors:
    Gray = "#5d5d5d"
    LightGray = "#fafafa"
    Black = "#000000"
    White = "#FFFFFF"
    Teal = "#008080"
    Aquamarine = "#76c8c8"
    Blue = "#2596be"
    LightCyan = "#badbdb"
    WhiteSmoke = "#dedad2"
    Cream = "#e4bcad"
    PeachPuff = "#df979e"
    HotPink = "#d7658b"
    DeepPink = "#c80064"
    LightSeaGreen = "#20B2AA"
    DarkGray = "#464144"

#RF
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmrf = confusion_matrix(y_test, y_predictedrf)
colors = ['lightgray', Colors.Teal, Colors.Teal, Colors.Teal, Colors.Teal,
      ↪ Colors.Teal, Colors.Teal, Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmrf, cmap=colormap, annot=True, fmt="d", linewidths=1, cbar=False,
            yticklabels=['Actual Non-Subscribed', 'Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed', 'Predicted_
      ↪ Subscribed'], annot_kws={"fontsize":14})
print("Random Forest", "\n", classification_report(y_test, y_predictedrf))
ax0.text(0, -0.75, 'Random Forest_
      ↪ Performance', fontsize=18, fontweight='bold', fontfamily='serif')

#DT
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
```

```

ax0 = fig.add_subplot(gs[0:1, 0:1])

cmdt = confusion_matrix(y_test,y_predicteddt)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmdt, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↪Subscribed'],annot_kws={"fontsize":14})
print("Decision Tree","\n",classification_report(y_test, y_predicteddt))
ax0.text(0,-0.75,'Decision Tree_
↪Performance',fontsize=18,fontweight='bold',fontfamily='serif')

#logistic regression
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmlgr = confusion_matrix(y_test,y_predictedlgr)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmlgr, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↪Subscribed'],annot_kws={"fontsize":14})
print("Logistic Regression","\n",classification_report(y_test, y_predictedlgr))
ax0.text(0,-0.75,'Logistic Regression_
↪Performance',fontsize=18,fontweight='bold',fontfamily='serif')

```

Random Forest

	precision	recall	f1-score	support
0	0.95	0.93	0.94	13276
1	0.55	0.63	0.59	1780
accuracy			0.90	15056
macro avg	0.75	0.78	0.76	15056
weighted avg	0.90	0.90	0.90	15056

Decision Tree

	precision	recall	f1-score	support
0	0.95	0.90	0.92	13276
1	0.45	0.64	0.53	1780



accuracy			0.87	15056
macro avg	0.70	0.77	0.73	15056
weighted avg	0.89	0.87	0.88	15056

Logistic Regression

	precision	recall	f1-score	support
0	0.95	0.92	0.94	13276
1	0.53	0.64	0.58	1780

accuracy			0.89	15056
macro avg	0.74	0.78	0.76	15056
weighted avg	0.90	0.89	0.89	15056

[39]: Text(0, -0.75, 'Logistic Regression Performance')



Logistic Regression Performance		
Actual Non-Subscribed	12240	1036
Actual Subscribed	634	1146
	Predicted Non-Subscribed	Predicted Subscribed

5.2.9 8. SMOTE- Calculate the AUROC

1. Higher is better
2. Between 0 and 1

```
[40]: #to get the probabilities of the prediction for each instance in the test set
np.random.seed(2)
rf_probs= rf.predict_proba(x_test)
dt_probs= dt.predict_proba(x_test)
lgr_probs= lgr.predict_proba(x_test)

#to keep the probabilities of positive outcomes
rf_probs=rf_probs[:,1]
rf_probs

dt_probs=dt_probs[:,1]
dt_probs

lgr_probs=lgr_probs[:,1]
lgr_probs

#print AUROC score
from sklearn.metrics import roc_curve, roc_auc_score
rf_auc=roc_auc_score(y_test,rf_probs)
dt_auc=roc_auc_score(y_test,dt_probs)
lgr_auc=roc_auc_score(y_test,lgr_probs)

print("Random Forest: AUROC= %.3f" % (rf_auc))
```

```

print("Decision Tree: AUROC= %.3f" % (dt_auc))
print("Logistic Regression: AUROC= %.3f" % (lgr_auc))

#Calculate and plot ROC curve
#roc_curve() returns three arrays, first one is false positive rate(fpr),
    ↳second one is true positive rate(tpr)
#and the third one is the threshold (_).
#For more information: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_curve.html
    ↳sklearn.metrics.roc_curve.html'''
rf_fpr,rf_tpr,_=roc_curve(y_test,rf_probs)
dt_fpr,dt_tpr,_=roc_curve(y_test,dt_probs)
lgr_fpr,lgr_tpr,_=roc_curve(y_test,lgr_probs)

#plot each model:
import matplotlib.pyplot as plt
plt.plot(rf_fpr,rf_tpr,marker=".",label="SMOTE- Random Forest (AUROC= %.3f)" % rf_auc)
    ↳rf_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

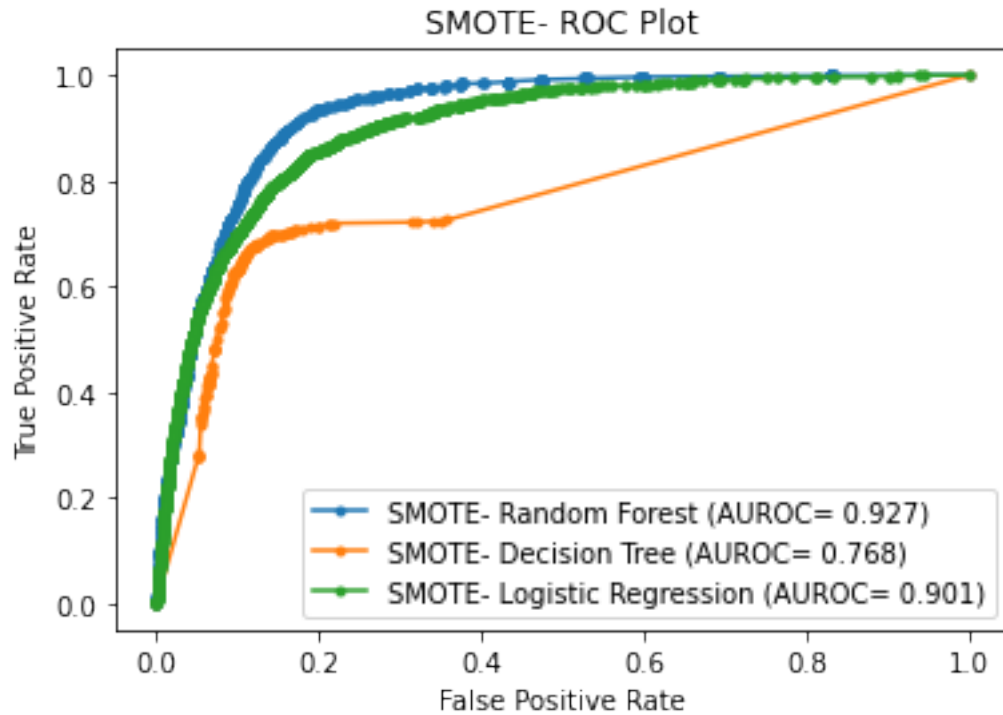
plt.plot(dt_fpr,dt_tpr,marker=".",label="SMOTE- Decision Tree (AUROC= %.3f)" % dt_auc)
    ↳dt_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.plot(lgr_fpr,lgr_tpr,marker=".",label="SMOTE- Logistic Regression (AUROC= %.3f)" % lgr_auc)
    ↳%.3f)" % lgr_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.title("SMOTE- ROC Plot")
plt.legend()
plt.show()
#The results show the best model AUROC is Random forest with 0.93 score
    ↳followed by Logisitic regression and Decision tree respectivilly.
#Comparing to the imbalanced section, the AUC has been decreased alittle bit as
    ↳
#logistic regression AUC decreased from 0.90 to 0.89 and Decision tree
    ↳decreased from 0.86 to 0.73.
#Random Forest AUC is stable before and after SMOTE.

```

Random Forest: AUROC= 0.927
 Decision Tree: AUROC= 0.768
 Logistic Regression: AUROC= 0.901



5.2.10 9. SMOTE- Summary (ACC,Recall, specificity,AUC,MCC, brier_score_loss,F1 Score, run time) for the models

```
[41]: np.random.seed(2)
from sklearn.metrics import roc_curve,
    ↳roc_auc_score,accuracy_score,recall_score,brier_score_loss,f1_score,matthews_corrcoef
from sklearn import tree, linear_model,ensemble
SmoteMLA = [tree.
    ↳DecisionTreeClassifier(max_depth=SMT_DT_max_depth,random_state=2),
        linear_model.
    ↳LogisticRegression(max_iter=SMT_lgr_max_iter,solver=SMT_lgr_solver,random_state=2),

        ensemble.RandomForestClassifier(n_estimators=SMT_RF_n_estimators,
    ↳max_depth=SMT_RF_max_depth,random_state=2)]
import time
name = []
Accuracy = []
Specificity=[]
Sensetivity=[]
alg_auc=[]
T_time=[]
P_time=[]
```

```

Mcc=[]
br_s=[]
f1score=[]
for alg in SmoteMLA:
    name.append(alg.__class__.__name__)

    if alg.__class__.__name__=="RandomForestClassifier":
        x_Train=rfx_train_res
        y_Train=rfy_train_res
    elif alg.__class__.__name__=="DecisionTreeClassifier":
        x_Train=dtx_train_res
        y_Train=dt_y_train_res
    elif alg.__class__.__name__=="LogisticRegression":
        x_Train=lgrx_train_res
        y_Train=lgry_train_res

    start=time.time()
    alg.fit(x_Train,y_Train)
    end=time.time()

    startp=time.time()
    y_predicted=alg.predict(x_test)
    endp=time.time()
    Accuracy.append(accuracy_score(y_test,y_predicted))
    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))
    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))
    Mcc.append(matthews_corrcoef(y_test, y_predicted))
    br_s.append(brier_score_loss(y_test, y_predicted))
    f1score.append(f1_score(y_test, y_predicted))
    alg_probs= alg.predict_proba(x_test)
    alg_probs=alg_probs[:,1]
    alg_auc.append(roc_auc_score(y_test,alg_probs))
    T_time.append(end-start)
    P_time.append(endp-startp)

Smotecomparison = pd.DataFrame({"name": name, "Accuracy":_
    ↳Accuracy,"Sensetivity":Sensetivity,
                                "Specificity":Specificity,"AUC":alg_auc,"MCC":
    ↳Mcc,
                                "Brier_score":br_s,"F1Score":
    ↳f1score,"TrainingTime_s": T_time,
                                "Prediction Time [s]": P_time})
display(Smotecomparison)

```



```
#From the below results we can conclude that the random forest is slightly
↳better than other models specially
#if we compared them by the Brier_score, Matthew's correlation coefficient
↳(MCC),Accuracy, and Specificity.

#MCC: Recently alot of researches prove that it is a better measurement
↳specially
#for binary classification problems along with the brier score.
#As we tuned the models and SMOTE to improve the F1 score. We can notice a
↳slight improvement in MCC, F1 Score, and Sensitivity.
```

	name	Accuracy	Sensetivity	Specificity	AUC \
0	DecisionTreeClassifier	0.866432	0.642135	0.896505	0.768310
1	LogisticRegression	0.889081	0.643820	0.921964	0.900733
2	RandomForestClassifier	0.895191	0.632022	0.930476	0.926821

	MCC	Brier_score	F1Score	TrainingTime_s	Prediction Time [s]
0	0.466089	0.133568	0.531999	0.682102	0.007563
1	0.518932	0.110919	0.578496	0.684723	0.007331
2	0.529779	0.104809	0.587774	11.087643	0.573028

5.2.11 10. SMOTE- Check Over/Underfitting

Brier Score was used to evaluate the model fitting

The reason behind using the Brier score metric is that Brier score measures the mean squared difference between the predicted probability and the actual outcome. The Brier score always takes on a value between zero and one, since this is the largest possible difference between a predicted probability (which must be between zero and one) and the actual outcome (which can take on values of only 0 and 1). It can be decomposed is the sum of refinement loss and calibration loss. The Brier score is appropriate for binary and categorical outcomes that can be structured as true or false, but is inappropriate for ordinal variables which can take on three or more values

Brier score is actually 1-Accuracy

”The Brier score is a strictly proper scoring function that is equivalent to the mean squared error” (Chicco, D., Warrens, M. J., & Jurman, G. (2021)).

*The smaller the Brier score loss, the better, hence the naming with “loss”.

1. If the performance of the model on the training set is poor that means the model is underfitting to the test data.
2. If the performance of the model on the test set is poor that means the model is overfitting to the training data.

```
[ ]: from sklearn.metrics import roc_curve,
↳roc_auc_score,accuracy_score,recall_score,brier_score_loss
```

```

from sklearn import tree, linear_model, ensemble
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import validation_curve
from sklearn.model_selection import KFold, RepeatedKFold
from sklearn.model_selection import cross_val_score
import numpy as np
np.random.seed(1)
import matplotlib.pyplot as plt
plt.style.use("ggplot")

```

```

[ ]: #Decision Tree
# prepare the cross-validation procedure
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
SMT_DT_list_training_BrierScore=[]
SMT_DT_list_testing_BrierScore=[]

X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
    model=tree.DecisionTreeClassifier(max_depth=SMT_DT_max_depth,random_state=2)

# apply the tuned smote to balance the training set:
smdt=SMOTE(sampling_strategy=DT_SMT_sampling_strategy,random_state=10)
dtx_train_res,dt_y_train_res=smdt.fit_resample(X_train,Y_train)

model.fit(dtx_train_res,dt_y_train_res)

y_train_data_pred=model.predict(dtx_train_res)
y_test_data_pred=model.predict(X_test)

↳
↳SMT_DT_fold_training_BrierScore=brier_score_loss(dt_y_train_res,y_train_data_pred)
SMT_DT_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

SMT_DT_list_training_BrierScore.append(SMT_DT_fold_training_BrierScore)
SMT_DT_list_testing_BrierScore.append(SMT_DT_fold_testing_BrierScore)

```

```

[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
↳phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(SMT_DT_list_training_BrierScore).
↳ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")

```

```
plt.title("Decision Tree \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(SMT_DT_list_testing_BrierScore).
        ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title(" Decision Tree \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()
```

#From the below graphs if we looked at the the first graph " Decision
↪Tree-Training BrierScore across folds" we can see that the model
#its really well as the BrierScore is around 0.03.
#In the second graph "Decision Tree -Testing BrierScore across folds"it also
↪fits well as the BrierScore on the test set is around 0.1 .
#So it fits the training and the test set.
#For sure random forest performs better as we can see that in the random forest
↪test below.



```
[ ]: #Random Forest
# prepare the cross-validation procedure
cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)
SMT_RF_list_training_BrierScore=[]
SMT_RF_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(dff):
    X_train,X_test=X.iloc[train_index:],X.iloc[test_index,:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
```

```

model=ensemble.RandomForestClassifier(n_estimators=SMT_RF_n_estimators,
↪max_depth=SMT_RF_max_depth,random_state=2)

# apply the tuned smote to balance the training set:
smrf=SMOTE(sampling_strategy=RF_SMT_sampling_strategy,random_state=10)
rfx_train_res,rfy_train_res=smrf.fit_resample(X_train,Y_train)

model.fit(rfx_train_res,rfy_train_res)

y_train_data_pred=model.predict(rfx_train_res)
y_test_data_pred=model.predict(X_test)

↪
↪SMT_RF_fold_training_BrierScore=brier_score_loss(rfy_train_res,y_train_data_pred)
SMT_RF_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

SMT_RF_list_training_BrierScore.append(SMT_RF_fold_training_BrierScore)
SMT_RF_list_testing_BrierScore.append(SMT_RF_fold_testing_BrierScore)


```

```

[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
↪phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(SMT_RF_list_training_BrierScore).
↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Random Forest \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(SMT_RF_list_testing_BrierScore).
↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title(" Random Forest \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

#From the below graphs if we looked at the the first graph "Random
↪Forest-Training BrierScore across folds"
#we can see the the models fits really well as the BrierScore is almost 0
#In the second graph "Random Forest -Testing BrierScore across folds"it also
↪fits well as the BrierScore on the test set is around 0.1

```



```
[ ]: #Logistic Regression
# prepare the cross-validation procedure

cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)
SMT_lgr_list_training_BrierScore=[]
SMT_lgr_list_testing_BrierScore=[]
X=dff.drop(["y"],axis="columns")
y=dff.y
for train_index,test_index in cv.split(dff):
    X_train,X_test=X.iloc[train_index:],X.iloc[test_index:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
    model=linear_model.
    ↳LogisticRegression(max_iter=SMT_lgr_max_iter,solver=SMT_lgr_solver,random_state=2)

# apply the tuned smote to balance the training set:
smlgr=SMOTE(sampling_strategy=lgr_SMT_sampling_strategy,random_state=10)
lgrx_train_res,lgry_train_res=smlgr.fit_resample(X_train,Y_train)

model.fit(lgrx_train_res,lgry_train_res)

y_train_data_pred=model.predict(lgrx_train_res)
y_test_data_pred=model.predict(X_test)

↳
↳SMT_lgr_fold_training_BrierScore=brier_score_loss(lgry_train_res,y_train_data_pred)
SMT_lgr_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

SMT_lgr_list_training_BrierScore.append(SMT_lgr_fold_training_BrierScore)
SMT_lgr_list_testing_BrierScore.append(SMT_lgr_fold_testing_BrierScore)
```

```
[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
      ↪phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.
      ↪array(SMT_lgr_list_training_BrierScore).ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Logistic Regression \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(SMT_lgr_list_testing_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title("Logistic Regression \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

#From the below graphs if we looked at the the first graph " Logistic
      ↪Regression-Training BrierScore across folds" we can see the the models fits
      ↪really well as the BrierScore is around 0.1
#In the second graph "Logistic Regression -Testing BrierScore across folds"it
      ↪also fits well as the BrierScore on the test set is around 0.1 .
#That means that we have a fitted model and stable one.
```



```
[ ]: display("LG_Test_BScore_Mean:", round(mean(SMT_lgr_list_testing_BrierScore),3))
display("LG_Train_BScore_Mean:",
      ↪round(mean(SMT_lgr_list_training_BrierScore),3))

display("DT_Test_BScore_Mean:", round(mean(SMT_DT_list_testing_BrierScore),3))
```

```
display("DT_Train_BScore_Mean:", round(mean(SMT_DT_list_training_BrierScore),3))

display("RF_Test_BScore_Mean:", round(mean(SMT_RF_list_testing_BrierScore),3))
display("RF_Train_BScore_Mean:", round(mean(SMT_RF_list_training_BrierScore),3))
```

'LG_Test_BScore_Mean: '

0.117

'LG_Train_BScore_Mean: '

0.105

'DT_Test_BScore_Mean: '

0.132

'DT_Train_BScore_Mean: '

0.035

'RF_Test_BScore_Mean: '

0.104

'RF_Train_BScore_Mean: '

0.0

5.2.12 11. SMOTE Summary without Tuning

```
[42]: from sklearn.metrics import roc_curve, \
      ↳roc_auc_score, accuracy_score, recall_score, brier_score_loss, matthews_corrcoef, f1_score
from sklearn import tree, linear_model, ensemble #svm, tree, linear_model, \
      ↳neighbors, naive_bayes, ensemble, discriminant_analysis, gaussian_process
from imblearn.over_sampling import SMOTE

SmoteMLA= [tree.DecisionTreeClassifier(random_state=14), #not tuned
          linear_model.LogisticRegression(random_state=15, n_jobs=-1), #not tuned,
```

```

        ensemble.RandomForestClassifier(random_state=13)]#not tuned]

sm= SMOTE(random_state=101)#SMOTE function without detremining and tuning the
↳sampling strategy

import time
name = []
Accuracy = []
Specificity=[]
Sensetivity=[]
alg_auc=[]
T_time=[]
P_time=[]
Mcc=[]
br_s=[]
f1score=[]

for alg in SmoteMLA:
    name.append(alg.__class__.__name__)

    #resampling
    x_Train_res,y_Train_res=sm.fit_resample(x_train,y_train)

    #train the models
    start=time.time()
    alg.fit(x_Train_res,y_Train_res)
    end=time.time()

    #prediting the test set by each model
    startp=time.time()
    y_predicted=alg.predict(x_test)
    endp=time.time()

    #calculating accuracy, specificity, sensetivity, MCC, brier score, F1 score
    Accuracy.append(accuracy_score(y_test,y_predicted))
    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))
    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))
    Mcc.append(matthews_corrcoef(y_test, y_predicted))
    br_s.append(brier_score_loss(y_test, y_predicted))
    f1score.append(f1_score(y_test, y_predicted))

    #computing ROC-AUC
    alg_probs= alg.predict_proba(x_test)
    alg_probs=alg_probs[:,1]
    alg_auc.append(roc_auc_score(y_test,alg_probs))

    #counting the running time of prediting and training mdoel

```



```

T_time.append(end-start)
P_time.append(endp-startp)

SMTcomparison_NotTuned = pd.DataFrame({"name": name, "Accuracy":_
↳Accuracy, "Sensetivity":Sensetivity,
"Specificity":Specificity, "AUC":
↳alg_auc, "Mcc":Mcc,
"brier_score_loss":br_s, "F1Score":
↳f1score,
"Training Time [s]": T_time, "Prediction_
↳Time [s]": P_time})
display(SMTcomparison_NotTuned)
#If we compared the below result with the results of Smotecomparison, we can_
↳notice that the tuning didnt show a notable improvement in the results.

```

	name	Accuracy	Sensetivity	Specificity	AUC \
0	DecisionTreeClassifier	0.860720	0.557303	0.901401	0.729352
1	LogisticRegression	0.890077	0.629213	0.925053	0.899705
2	RandomForestClassifier	0.894660	0.661236	0.925957	0.926318

	Mcc	brier_score_loss	F1Score	Training Time [s] \
0	0.411602	0.139280	0.486155	0.817559
1	0.515019	0.109923	0.575096	2.168229
2	0.540841	0.105340	0.597462	10.280462

	Prediction Time [s]
0	0.008777
1	0.006218
2	0.477386

5.3 Random Undersampling (RUS)

1. In this stage the models have been built on balanced dataset using Random Undersampling.
2. Sampling strategy parameter in Undersampling function has been tuned for each model in order to optimize the F1-score.
3. The models have been tuned based on some parameters: A. Random Forest tuned parameters are: Max depth and the n_estimators (number of trees) B. Decision Tree tuned parameters are: Max depth. C. Logistic Regression tuned parameters are : Max iteration and Solver.
4. We focused on optimizing F1-Score metric as it was the smallest value in the confusion matrix in the first run

If the performance is similar or better to SMOTE then using RUS would be chosen as this method would save more time to train models and predict outputs.

5.3.1 1. RUS Hyperparameter Tuning (based on Sampling strategy and F1-Score)

```
[ ]: #Tuning randomundersampler for Random Forest
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from imblearn.under_sampling import RandomUnderSampler
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
import imblearn

[ ]: #Tuning randomundersampler for Random Forest
rfweights=np.linspace(0.3,0.5,50)
pipe=imblearn.pipeline.
    ↳make_pipeline(RandomUnderSampler(),RandomForestClassifier())
gsc=GridSearchCV(estimator=pipe,param_grid={"randomundersampler__sampling_strategy":
    ↳rfweights},scoring="f1",cv=3,n_jobs=-1)
RUSrf_grid_result=gsc.fit(x_train,y_train)

#Tuning randomundersampler for Decision Tree
dtweights=np.linspace(0.8,1,50)
pipe=imblearn.pipeline.
    ↳make_pipeline(RandomUnderSampler(),DecisionTreeClassifier())
gsc=GridSearchCV(estimator=pipe,param_grid={"randomundersampler__sampling_strategy":
    ↳dtweights},scoring="f1",cv=3,n_jobs=-1)
RUSdt_grid_result=gsc.fit(x_train,y_train)

#Tuning randomundersampler for Logistic Regression
lgrweights=np.linspace(0.4,1,20)
pipe=imblearn.pipeline.
    ↳make_pipeline(RandomUnderSampler(),LogisticRegression(solver="liblinear"))
gsc=GridSearchCV(estimator=pipe,param_grid={"randomundersampler__sampling_strategy":
    ↳lgrweights},scoring="f1",cv=3,n_jobs=-1)
RUSlgr_grid_result=gsc.fit(x_train,y_train)

[ ]: #RF
print("RF-the best parameters are %s with a score of %0.2f" %_
    ↳(RUSrf_grid_result.best_params_,RUSrf_grid_result.best_score_))#Mean_
    ↳cross-validated score of the best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":RUSrf_grid_result.
    ↳cv_results_["mean_test_score"],"weight":rfweights})
#the score is the average of all cv folds for a single combination of the_
    ↳parameters you specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Random Forest")

#DT
```

```

print("DT-the best parameters are %s with a score of %0.2f" %
      (RUSdt_grid_result.best_params_,RUSdt_grid_result.best_score_))#Mean
      ↳cross-validated score of the best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":RUSdt_grid_result.
      ↳cv_results_["mean_test_score"],"weight":dtweights}) #the score is the
      ↳average of all cv folds for a single combination of the parameters you
      ↳specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Decision Tree")

#Logistic Regression
print("LGR-the best parameters are %s with a score of %0.2f" %
      (RUSlgr_grid_result.best_params_,RUSlgr_grid_result.best_score_))#Mean
      ↳cross-validated score of the best_estimator
weight_roc_auc_score_df=pd.DataFrame({"score":RUSlgr_grid_result.
      ↳cv_results_["mean_test_score"],"weight":lgrweights}) #the score is the
      ↳average of all cv folds for a single combination of the parameters you
      ↳specify in the tuned_params
weight_roc_auc_score_df.plot(x="weight",title="Logistic Regression")



```


RF-the best parameters are {'randomundersampler__sampling_strategy':
0.3816326530612245} with a score of 0.61

DT-the best parameters are {'randomundersampler__sampling_strategy': 0.8} with a
score of 0.49

LGR-the best parameters are {'randomundersampler__sampling_strategy':
0.5578947368421052} with a score of 0.57

```
[ ]: <AxesSubplot: title={'center': 'Logistic Regression'}, xlabel='weight'>
```





5.3.2 2. Data Balancing: (Random UnderSampling)

```
[43]: #Tuning randomundersampler for Random Forest
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from imblearn.under_sampling import RandomUnderSampler
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
import imblearn
```

```
[44]: #from the above we can conclude that the best parameters for Random
      ↳ Undersampling are:

      #Please note that the sampling strategy is the desired ratio of the number of
      ↳ samples in the minority
      #class over the number of samples in the majority class after resampling.'''

      #Random Forest:
      RF_RUS_sampling_strategy=0.3816326530612245

      # Decision Tree:
      DT_RUS_sampling_strategy=0.8
```

```
# Logistic Regression:
lgr_RUS_sampling_strategy=0.5578947368421052

#These values may differ slightly on each run of the previous tunning.
#Please note that I have saved the random undersampling -sampling strategy
↳ values (after tuning) for each model
#in the above variables in order to make the run time shorter everytime we need
↳ to balance the data.
```

```
[45]: #The portions for minority and majority before RUS
import numpy as np
unique,count=np.unique(y_train,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[45]: {0: 26646, 1: 3509}
```

```
[46]: #Resample DT
RUSdt=RandomUnderSampler(sampling_strategy=DT_RUS_sampling_strategy,random_state=100)
dtx_train_rus,dt_y_train_rus=RUSdt.fit_resample(x_train,y_train)
unique,count=np.unique(dt_y_train_rus,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[46]: {0: 4386, 1: 3509}
```

```
[47]: #Resample RF
RUSrf=RandomUnderSampler(sampling_strategy=RF_RUS_sampling_strategy,random_state=11)#
rfx_train_rus,rfy_train_rus=RUSrf.fit_resample(x_train,y_train)
unique,count=np.unique(rfy_train_rus,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[47]: {0: 9194, 1: 3509}
```

```
[48]: #Resample logistic regression
RUSlgr=RandomUnderSampler(sampling_strategy=lgr_RUS_sampling_strategy,random_state=12)
lgrx_train_rus,lgry_train_rus=RUSlgr.fit_resample(x_train,y_train)
unique,count=np.unique(lgry_train_rus,return_counts=True)
y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
y_train_dict_value_count
```

```
[48]: {0: 6289, 1: 3509}
```

5.3.3 3.RUS- Random Forest Hyperparameter Tuning (based on max depth, number of trees and F1-Score)

```
[ ]: from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(random_state=3)

[ ]: import numpy as np
from sklearn.model_selection import GridSearchCV

max_depth_range=np.arange(14,20,1)
n_estimators_range=np.arange(10,110,10)
param_grid=dict(max_depth=max_depth_range,n_estimators=n_estimators_range)
rfgridRUS=GridSearchCV(estimator=rf,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)

[ ]: rfgridRUS.fit(rfx_train_rus,rfy_train_rus)

[ ]: GridSearchCV(cv=3, estimator=RandomForestClassifier(random_state=3), n_jobs=-1,
                param_grid={'max_depth': array([14, 15, 16, 17, 18, 19]),
                             'n_estimators': array([ 10, 20, 30, 40, 50, 60,
70, 80, 90, 100])},
                scoring='f1')

[ ]: print("the best parameters are %s with a score of %0.2f" % (rfgridRUS.
    ↳best_params_,rfgridRUS.best_score_))#Mean cross-validated score of the
    ↳best_estimator
#the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params
```

the best parameters are {'max_depth': 17, 'n_estimators': 80} with a score of 0.73

```
[ ]: #data frame of grid search parameters and their accuracy scores
import pandas as pd

grid_results = pd.concat([pd.DataFrame(rfgridRUS.cv_results_["params"]),
                          pd.DataFrame(rfgridRUS.
    ↳cv_results_["mean_test_score"], columns=["f1"])],axis=1)
grid_results.head()

#preparing data for making contour plots
grid_contour = grid_results.groupby(['max_depth','n_estimators']).mean()
grid_contour

#pivot data:
grid_reset=grid_contour.reset_index()
grid_reset.columns=["max_depth","n_estimator","f1"]
grid_pivot=grid_reset.pivot("max_depth","n_estimator")
```

```

#assigning the pivoted data into the respective x,y, and z variables
x=grid_pivot.columns.levels[1].values
y=grid_pivot.index.values
z=grid_pivot.values

#2D contour plot
import plotly.graph_objects as go

# X and Y axes labels
layout = go.Layout(
    xaxis=go.layout.XAxis(
        title=go.layout.xaxis.Title(
            text='n_estimators')
        ),
    yaxis=go.layout.YAxis(
        title=go.layout.yaxis.Title(
            text='max_depth')
    ) )

fig = go.Figure(data = [go.Contour(z=z, x=x, y=y)], layout=layout )

fig.update_layout(title='Hyperparameter tuning', autosize=False,
                    width=500, height=500,
                    margin=dict(l=65, r=50, b=65, t=90))

fig.show()

```

5.3.4 4. RUS- Decision Tree Hyperparameter Tuning (based on max depth and F1-Score)

```
[ ]: from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(random_state=3)
```

```
[ ]: import numpy as np
from sklearn.model_selection import GridSearchCV

max_depth_range=np.arange(1,14,2)
param_grid=dict(max_depth=max_depth_range)
dtgridRUS=GridSearchCV(estimator=dt,param_grid=param_grid,scoring="roc_auc",cv=10,n_jobs=-1)
```

```
[ ]: np.arange(1,14,2)
```

```
[ ]: array([ 1,  3,  5,  7,  9, 11, 13])
```

```
[ ]: dtgridRUS.fit(dtx_train_rus,dy_train_rus)
```



```
[ ]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(random_state=3), n_jobs=-1,
                  param_grid={'max_depth': array([ 1,  3,  5,  7,  9, 11, 13])},
                  scoring='roc_auc')
```

```
[ ]: print("the best parameters are %s with a score of %0.2f" % (dtgridRUS.
    ↳best_params_, dtgridRUS.best_score_)) #Mean cross-validated score of the
    ↳best_estimator
    #the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params
```

the best parameters are {'max_depth': 7} with a score of 0.87

5.3.5 3.RUS- Logistic Regression Hyperparameter Tuning (based on max Iteration, Solver, and F1-Score)

```
[ ]: from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(random_state=3,n_jobs=-1)
```

```
[ ]: #This step has been run repeatedly on different ranges of the model parameters
    ↳in order to tune them on the right values.
import numpy as np
from sklearn.model_selection import GridSearchCV

max_iter=np.arange(60,100,1)
param_grid=dict(max_iter=max_iter,solver=['liblinear', 'lbfgs'])

lgrgridRUS=GridSearchCV(estimator=lgr,param_grid=param_grid,scoring="f1",cv=3,n_jobs=-1)
```

```
[ ]: lgrgridRUS.fit(rfx_train_rus,rfy_train_rus)
```

```
[ ]: GridSearchCV(cv=3, estimator=LogisticRegression(n_jobs=-1, random_state=3),
                  n_jobs=-1,
                  param_grid={'max_iter': array([60, 61, 62, 63, 64, 65, 66, 67, 68,
69, 70, 71, 72, 73, 74, 75, 76,
77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93,
94, 95, 96, 97, 98, 99])},
                  'solver': ['liblinear', 'lbfgs']},
                  scoring='f1')
```

```
[ ]: print("the best parameters are %s with a score of %0.2f" % (lgrgridRUS.
    ↳best_params_, lgrgridRUS.best_score_)) #Mean cross-validated score of the
    ↳best_estimator
    #the score is the average of all cv folds for a single combination of the
    ↳parameters you specify in the tuned_params
    #lbfgs stayed as a better choice.
```

the best parameters are {'max_iter': 69, 'solver': 'lbfgs'} with a score of 0.65

5.3.6 5. RUS- Train the model and calculate the run time of training the model

```
[49]: #from the above we can conclude that the best parameters are:

#Random Forest:
RUS_RF_max_depth= 17
RUS_RF_n_estimators= 80

# Decision Tree:

RUS_DT_max_depth=7

#logistic Regression:
RUS_lgr_max_iter=69
RUS_lgr_solver="lbfgs"

#These values may differ slightly on each run of the previous tunning.
#Please note that I have saved the parameters values (after tuning) of each
→model in the above variables
#in order to make the run time shorter everytime we need to train the model.

[50]: #RF
import time
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=RUS_RF_n_estimators, max_depth=
→RUS_RF_max_depth,random_state=3)

startrf=time.time()
rf.fit(rfx_train_rus,rfy_train_rus)
endrf=time.time()

#DT
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(max_depth=RUS_DT_max_depth,random_state=3)
startdt=time.time()
dt.fit(dtx_train_rus,dy_train_rus)
enddt=time.time()

#Logistic Regression
from sklearn.linear_model import LogisticRegression
lgr=LogisticRegression(max_iter=RUS_lgr_max_iter,solver=RUS_lgr_solver,n_jobs=-1)
# -1 means using all processors. This means that the n_jobs parameter can be
→used to distribute and exploit all the CPUs available in the local computer
startlgr=time.time()
```

```

lgr.fit(lgrx_train_rus,lgry_train_rus)
endlgr=time.time()

print("RF Training Run Time [s]:",endrf-startrf)
print("DT Training Run Time [s]:",enddt-startdt)
print("LGR Training Run Time [s]:",endlgr-startlgr)
#The results shows the the faster model is the Decision tree,
#Logistic regression, and the Random forest respectively.
#Random Forest model was the slowest one( which is expected as it is an
→ ensemble model)

```

```

RF Training Run Time [s]: 1.3602180480957031
DT Training Run Time [s]: 0.04312729835510254
LGR Training Run Time [s]: 1.021684169769287

```

5.3.7 6. RUS- Run Time for prediction

```

[51]: # the run time to predict the test set
startrf=time.time()
y_predictedrf=rf.predict(x_test)
endrf=time.time()

startdt=time.time()
y_predicteiddt=dt.predict(x_test)
enddt=time.time()

startlgr=time.time()
y_predictedlgr=lgr.predict(x_test)
endlgr=time.time()

print("RF Run Time to Predict the testset [s]:",endrf-startrf)
print("DT Run Time to Predict the testset [s]:",enddt-startdt)
print("LGR Run Time to Predict the testset [s]:",endlgr-startlgr)

# The results shows the the faster model is the Logistic regression,Decision
→ tree, and the Random forest respectively.
# Random Forest model was the slowest one( which is expected as it is an
→ ensemble model)

```

```

RF Run Time to Predict the testset [s]: 0.2660255432128906
DT Run Time to Predict the testset [s]: 0.006936311721801758
LGR Run Time to Predict the testset [s]: 0.005370140075683594

```

5.3.8 7. RUS- Confusion Matrix

```
[52]: from sklearn.metrics import recall_score, make_scorer, confusion_matrix, \
      ↪ classification_report, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns

class Colors:
    Gray = "#5d5d5d"
    LightGray = "#fafafa"
    Black = "#000000"
    White = "#FFFFFF"
    Teal = "#008080"
    Aquamarine = "#76c8c8"
    Blue = "#2596be"
    LightCyan = "#badbdb"
    WhiteSmoke = "#dedad2"
    Cream = "#e4bcad"
    PeachPuff = "#df979e"
    HotPink = "#d7658b"
    DeepPink = "#c80064"
    LightSeaGreen = "#20B2AA"
    DarkGray = "#464144"

#RF
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmrf = confusion_matrix(y_test, y_predictedrf)
colors = ['lightgray', Colors.Teal, Colors.Teal, Colors.Teal, Colors.Teal, Colors.
      ↪ Teal, Colors.Teal, Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmrf, cmap=colormap, annot=True, fmt="d", linewidths=1, cbar=False,
            yticklabels=['Actual Non-Subscribed', 'Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed', 'Predicted_
      ↪ Subscribed'], annot_kws={"fontsize":14})
print("Random Forest", "\n", classification_report(y_test, y_predictedrf))
ax0.text(0, -0.75, 'Random Forest_
      ↪ Performance', fontsize=18, fontweight='bold', fontfamily='serif')

#DT
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])
```

```

cmdt = confusion_matrix(y_test,y_predicteddt)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmdt, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↪Subscribed'],annot_kws={"fontsize":14})
print("Decision Tree","\n",classification_report(y_test, y_predicteddt))
ax0.text(0,-0.75,'Decision Tree_
↪Performance',fontsize=18,fontweight='bold',fontfamily='serif')

#logistic regression
fig = plt.figure(figsize=(6, 2), dpi=100, facecolor=Colors.LightGray)
gs = fig.add_gridspec(1, 1)
gs.update(wspace=0.35, hspace=0.27)
ax0 = fig.add_subplot(gs[0:1, 0:1])

cmlgr = confusion_matrix(y_test,y_predictedlgr)
colors = ['lightgray',Colors.Teal,Colors.Teal,Colors.Teal,Colors.Teal,Colors.
↪Teal,Colors.Teal,Colors.Teal]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
sns.heatmap(cmlgr, cmap=colormap,annot=True,fmt="d", linewidths=1,cbar=False,
            yticklabels=['Actual Non-Subscribed','Actual Subscribed'],
            xticklabels=['Predicted Non-Subscribed','Predicted_
↪Subscribed'],annot_kws={"fontsize":14})
print("Logistic Regression","\n",classification_report(y_test, y_predictedlgr))
ax0.text(0,-0.75,'Logistic Regression_
↪Performance',fontsize=18,fontweight='bold',fontfamily='serif')

```

Random Forest

	precision	recall	f1-score	support
0	0.96	0.92	0.94	13276
1	0.54	0.72	0.62	1780
accuracy			0.89	15056
macro avg	0.75	0.82	0.78	15056
weighted avg	0.91	0.89	0.90	15056

Decision Tree

	precision	recall	f1-score	support
0	0.97	0.83	0.89	13276
1	0.39	0.81	0.53	1780



accuracy			0.83	15056
macro avg	0.68	0.82	0.71	15056
weighted avg	0.90	0.83	0.85	15056

Logistic Regression

	precision	recall	f1-score	support
0	0.95	0.91	0.93	13276
1	0.51	0.68	0.58	1780

accuracy			0.89	15056
macro avg	0.73	0.80	0.76	15056
weighted avg	0.90	0.89	0.89	15056

[52]: Text(0, -0.75, 'Logistic Regression Performance')



Logistic Regression Performance		
Actual Non-Subscribed	12131	1145
Actual Subscribed	575	1205
	Predicted Non-Subscribed	Predicted Subscribed

5.3.9 8.RUS- Calculate the AUROC

[53]: *#to get the probabilities of the prediction for each instance in the test set*

```

rf_probs= rf.predict_proba(x_test)
dt_probs= dt.predict_proba(x_test)
lgr_probs= lgr.predict_proba(x_test)

#to keep the probabilities of positive outcomes
rf_probs=rf_probs[:,1]
rf_probs

dt_probs=dt_probs[:,1]
dt_probs

lgr_probs=lgr_probs[:,1]
lgr_probs

#print AUROC score
from sklearn.metrics import roc_curve, roc_auc_score
rf_auc=roc_auc_score(y_test,rf_probs)
dt_auc=roc_auc_score(y_test,dt_probs)
lgr_auc=roc_auc_score(y_test,lgr_probs)

print("Random Forest: AUROC= %.3f" % (rf_auc))
print("Decision Tree: AUROC= %.3f" % (dt_auc))
print("Logistic Regression: AUROC= %.3f" % (lgr_auc))

```

```

#Calculate and plot ROC curve
#roc_curve() returns three arrays, first one is false positive rate(fpr),
#second one is true positive rate(tpr) and the third one is the threshold (_).
#For more information: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_curve.html
    ↳sklearn.metrics.roc_curve.html

rf_fpr,rf_tpr,_=roc_curve(y_test,rf_probs)
dt_fpr,dt_tpr,_=roc_curve(y_test,dt_probs)
lgr_fpr,lgr_tpr,_=roc_curve(y_test,lgr_probs)

import matplotlib.pyplot as plt
plt.plot(rf_fpr,rf_tpr,marker=".",label="RUS- Random Forest (AUROC= %0.3f)" %
    ↳rf_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.plot(dt_fpr,dt_tpr,marker=".",label="RUS- Decision Tree (AUROC= %0.3f)" %
    ↳dt_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.plot(lgr_fpr,lgr_tpr,marker=".",label="RUS- Logistic Regression (AUROC= %0.
    ↳3f)" % lgr_auc)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')


plt.title("Random Undersampling - ROC Plot")
plt.legend()
plt.show()

#The results show the best model AUROC is Random forest with 0.93 score
    ↳followed by
#Logisitic regression and Decision tree respectively.
#Comparing to the imbalanced section, the AUC has been increased slightly in
    ↳Decision tree from 0.84 to 0.86.
#Random Forest and logistic regression are stable before and after Random
    ↳UnderSampling.

#Comparing to the SMOTE section, the AUC has been increased slightly as
    ↳logistic regression AUC
# increased from 0.89 to 0.91 and Decision tree increased from 0.77 to 0.87.
#Random Forest is stable before and after Random UnderSampling.

```

Random Forest: AUROC= 0.929
 Decision Tree: AUROC= 0.870
 Logistic Regression: AUROC= 0.910



5.3.10 9. RUS- Summary (ACC,Recall, specificity,AUC,MCC, brier_score_loss, F1 Score, run time) for the models

```
[54]: from sklearn.metrics import roc_curve, \
    ↪roc_auc_score, accuracy_score, recall_score, brier_score_loss, matthews_corrcoef, f1_score
from sklearn import tree, linear_model, ensemble

RUSMLA = [tree.
    ↪DecisionTreeClassifier(max_depth=RUS_DT_max_depth, random_state=3),
    linear_model.
    ↪LogisticRegression(max_iter=RUS_lgr_max_iter, solver=RUS_lgr_solver, random_state=3, n_jobs=-1
        # -1 means using all processors. This means that the n_jobs parameter
    ↪can be used to distribute and exploit all the CPUs available in the local
    ↪computer
        ensemble.RandomForestClassifier(n_estimators=RUS_RF_n_estimators, \
    ↪max_depth=RUS_RF_max_depth, random_state=3)
    ]

import time
name = []
Accuracy = []
Specificity= []
```

```

Sensetivity=[]
alg_auc=[]
T_time=[]
P_time=[]
Mcc=[]
br_s=[]
f1score=[]
for alg in RUSMLA:
    name.append(alg.__class__.__name__)

    if alg.__class__.__name__=="RandomForestClassifier":
        x_Train=rfx_train_rus
        y_Train=rfy_train_rus
    elif alg.__class__.__name__=="DecisionTreeClassifier":
        x_Train=dtx_train_rus
        y_Train=dyt_train_rus
    elif alg.__class__.__name__=="LogisticRegression":
        x_Train=lgrx_train_rus
        y_Train=lgry_train_rus

    start=time.time()
    alg.fit(x_Train,y_Train)
    end=time.time()

    startp=time.time()
    y_predicted=alg.predict(x_test)
    endp=time.time()
    Accuracy.append(accuracy_score(y_test,y_predicted))
    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))
    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))
    Mcc.append(matthews_corrcoef(y_test, y_predicted))
    br_s.append(brier_score_loss(y_test, y_predicted))
    f1score.append(f1_score(y_test, y_predicted))
    alg_probs= alg.predict_proba(x_test)
    alg_probs=alg_probs[:,1]
    alg_auc.append(roc_auc_score(y_test,alg_probs))
    T_time.append(end-start)
    P_time.append(endp-startp)

RUScomparison = pd.DataFrame({"name": name, "Accuracy": Accuracy,"Sensetivity":
    ↳Sensetivity,"Specificity":Specificity,
                                "AUC":alg_auc,"Mcc":Mcc,"brier_score_loss":
    ↳br_s,"F1Score":f1score,

```

```

                                "TrainingTime_s": T_time, "Prediction Time [s]":
    ↪P_time})
display(RUScomparison)

#From the below results we can conclude that the random forest is slightly
    ↪better than other models
#specially if we compared them by the Brier_score,
#Matthew's correlation coefficient (MCC), Specificity, and AUC.
#Recently alot of researches prove that MCC is a better measurement specially
#for binary classification problems along with the brier score.
# we can notice if we compared the results of this summary with the summary of
    ↪SMOTE that the undersampling techniques improves the performance more than
    ↪applying SMOTE technique

```

	name	Accuracy	Sensetivity	Specificity	AUC \
0	DecisionTreeClassifier	0.827179	0.813483	0.829015	0.869632
1	LogisticRegression	0.885760	0.676966	0.913754	0.909707
2	RandomForestClassifier	0.894062	0.719101	0.917520	0.929225

	Mcc	brier_score_loss	F1Score	TrainingTime_s	Prediction Time [s]
0	0.481052	0.172821	0.526737	0.041147	0.007113
1	0.525518	0.114240	0.583535	0.609840	0.005797
2	0.563918	0.105938	0.616125	1.350638	0.287884

5.3.11 10. RUS- Check Over/Underfitting

Brier Score was used to evaluate the model fitting

The reason behind using the Brier score metric is that Brier score measures the mean squared difference between the predicted probability and the actual outcome. The Brier score always takes on a value between zero and one, since this is the largest possible difference between a predicted probability (which must be between zero and one) and the actual outcome (which can take on values of only 0 and 1). It can be decomposed is the sum of refinement loss and calibration loss. The Brier score is appropriate for binary and categorical outcomes that can be structured as true or false, but is inappropriate for ordinal variables which can take on three or more values

Brier score is actually 1-Accuracy

”The Brier score is a strictly proper scoring function that is equivalent to the mean squared error” (Chicco, D., Warrens, M. J., & Jurman, G. (2021)).

*The smaller the Brier score loss, the better, hence the naming with “loss”.

If the performance of the model on the training set is poor that means the model is underfitting to the test data.

If the performance of the model on the test set is poor that means the model is overfitting to the training data.

```
[ ]: from sklearn.metrics import roc_curve,
      ↪roc_auc_score, accuracy_score, recall_score, brier_score_loss
from sklearn import tree, linear_model, ensemble
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import validation_curve
from sklearn.model_selection import RepeatedKFold
import numpy as np
np.random.seed(3)
import matplotlib.pyplot as plt
plt.style.use("ggplot")
```

```
[ ]: #Decision Tree
# prepare the cross-validation procedure
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
RUS_DT_list_training_BrierScore=[]
RUS_DT_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
    model=tree.DecisionTreeClassifier(max_depth=RUS_DT_max_depth,random_state=3)

# apply the tuned Random Undersampling to balance the training set:
    ↪
    ↪RUSdt=RandomUnderSampler(sampling_strategy=DT_RUS_sampling_strategy,random_state=11)
    dtx_train_rus,dt_y_train_rus=RUSdt.fit_resample(X_train,Y_train)

    model.fit(dtx_train_rus,dt_y_train_rus)

    y_train_data_pred=model.predict(dtx_train_rus)
    y_test_data_pred=model.predict(X_test)

    ↪
    ↪RUS_DT_fold_training_BrierScore=brier_score_loss(dt_y_train_rus,y_train_data_pred)
    RUS_DT_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

    RUS_DT_list_training_BrierScore.append(RUS_DT_fold_training_BrierScore)
    RUS_DT_list_testing_BrierScore.append(RUS_DT_fold_testing_BrierScore)
```

```
[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
      ↪phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(RUS_DT_list_training_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
```

```

plt.ylabel("Training BrierScore")
plt.title("Decision Tree \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(RUS_DT_list_testing_BrierScore).
        ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title(" Decision Tree \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

```

#From the below graphs if we looked at the the first graph " Decision_
↪Tree-Training BrierScore across folds",
#we can see the the models fits really well as the BrierScore is around 0.2
#In the second graph "Decision Tree -Testing BrierScore across folds"it also_
↪fits well as the BrierScore on the test set is around 0.2.
#That means also that we have a fitted model and stable one.



```

[ ]: #Logistic Regression
# prepare the cross-validation procedure
cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)

RUS_lgr_list_training_BrierScore=[]
RUS_lgr_list_testing_BrierScore=[]

X=df.drop(["y"],axis="columns")
y=df.y

for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]

```

```

Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
model=linear_model.
↳ LogisticRegression(max_iter=RUS_lgr_max_iter,solver=RUS_lgr_solver,random_state=3,n_jobs=-1)

# apply the tuned RUS to balance/resample the training set:
↳
↳ RUSlgr=RandomUnderSampler(sampling_strategy=lgr_RUS_sampling_strategy,random_state=11)
lgrx_train_rus,lgry_train_rus=RUSlgr.fit_resample(X_train,Y_train)

model.fit(lgrx_train_rus,lgry_train_rus)

y_train_data_pred=model.predict(lgrx_train_rus)
y_test_data_pred=model.predict(X_test)

↳
↳ RUS_lgr_fold_training_BrierScore=brier_score_loss(lgry_train_rus,y_train_data_pred)
RUS_lgr_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

RUS_lgr_list_training_BrierScore.append(RUS_lgr_fold_training_BrierScore)
RUS_lgr_list_testing_BrierScore.append(RUS_lgr_fold_testing_BrierScore)

```

```

[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
↳ phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.
↳ array(RUS_lgr_list_training_BrierScore).ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Logistic Regression \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(RUS_lgr_list_testing_BrierScore).
↳ ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title("Logistic Regression \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

#From the below graphs if we looked at the the first graph "Logistic
↳ Regression-Training BrierScore across folds",
#we can see the the models fits well as the BrierScore is about 0.2
#In the second graph "Logistic Regression -Training BrierScore across folds"
#it also fits well as the BrierScore on the test set is about 0.1. that means
↳ also that we have a fitted model and stable one.

```



```
[ ]: #Random Forest
# prepare the cross-validation procedure
cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)
RUS_RF_list_training_BrierScore=[]
RUS_RF_list_testing_BrierScore=[]
X=df.drop(["y"],axis="columns")
y=df.y
for train_index,test_index in cv.split(df):
    X_train,X_test=X.iloc[train_index:],X.iloc[test_index:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]
    model=ensemble.RandomForestClassifier(n_estimators=RUS_RF_n_estimators,
    ↪max_depth=RUS_RF_max_depth,random_state=3)

# apply the tuned RUS to balance the training set:
↪RUSrf=RandomUnderSampler(sampling_strategy=RF_RUS_sampling_strategy,random_state=11)
rfx_train_rus,rfy_train_rus=RUSrf.fit_resample(X_train,Y_train)

model.fit(rfx_train_rus,rfy_train_rus)

y_train_data_pred=model.predict(rfx_train_rus)
y_test_data_pred=model.predict(X_test)

↪RUS_RF_fold_training_BrierScore=brier_score_loss(rfy_train_rus,y_train_data_pred)
RUS_RF_fold_testing_BrierScore=brier_score_loss(Y_test,y_test_data_pred)

RUS_RF_list_training_BrierScore.append(RUS_RF_fold_training_BrierScore)
RUS_RF_list_testing_BrierScore.append(RUS_RF_fold_testing_BrierScore)
```

```
[ ]: #plot the BrierScore of the training phase and the BrierScore of the testing
      ↪phase
plt.figure(figsize = (10, 4))
plt.subplot(1,2,1)
plt.plot(range(1,cv.get_n_splits()+1),np.array(RUS_RF_list_training_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Training BrierScore")
plt.title("Random Forest \n Training BrierScore across folds ")
plt.tight_layout()
plt.subplot(1,2,2)
plt.plot(range(1,cv.get_n_splits()+1),np.array(RUS_RF_list_testing_BrierScore).
      ↪ravel(),"o-")
plt.xlabel("Number of fold")
plt.ylabel("Testing BrierScore")
plt.title(" Random Forest \n Testing BrierScore across folds")
plt.tight_layout()
plt.show()

#From the below graphs if we looked at the the first graph "Random
      ↪Forest-Training BrierScore across folds"
#we can see the the models fits really well as the BrierScore is around 0
#in the second graph "Random Forest -Testing BrierScore across folds"it also
      ↪fits well as the BrierScore on the test set is around 0.1.
#That means also that we have a fitted model and stable one.
```



```
[ ]: display("LG_Test_BScore_Mean:", round(mean(RUS_lgr_list_testing_BrierScore),3))
display("LG_Train_BScore_Mean:",
      ↪round(mean(RUS_lgr_list_training_BrierScore),3))

display("DT_Test_BScore_Mean:", round(mean(RUS_DT_list_testing_BrierScore),3))
```



```
display("DT_Train_BScore_Mean:", round(mean(RUS_DT_list_training_BrierScore),3))

display("RF_Test_BScore_Mean:", round(mean(RUS_RF_list_testing_BrierScore),3))
display("RF_Train_BScore_Mean:", round(mean(RUS_RF_list_training_BrierScore),3))
```

```

      □
↳ -----

NameError                                Traceback (most recent call↳
↳ last)

    <ipython-input-67-a9236a495d38> in <module>
    ----> 1 display("LG_Test_BScore_Mean:",↳
↳ round(mean(RUS_lgr_list_testing_BrierScore),3))
          2 display("LG_Train_BScore_Mean:",↳
↳ round(mean(RUS_lgr_list_training_BrierScore),3))
          3
          4 display("DT_Test_BScore_Mean:",↳
↳ round(mean(RUS_DT_list_testing_BrierScore),3))
          5 display("DT_Train_BScore_Mean:",↳
↳ round(mean(RUS_DT_list_training_BrierScore),3))

NameError: name 'mean' is not defined
```

5.3.12 11. *RUS Summary without Tuning*

```
[55]: from sklearn.metrics import roc_curve,↳
↳ roc_auc_score, accuracy_score, recall_score, brier_score_loss, matthews_corrcoef, f1_score
from sklearn import tree, linear_model, ensemble

RUSMLA = [tree.DecisionTreeClassifier(random_state=4), #not tuned
linear_model.LogisticRegression(n_jobs=-1), #not tuned,
ensemble.RandomForestClassifier(random_state=3)] #not tuned]

RUS=RandomUnderSampler(random_state=100)
#random undersampling function without detremining and tuning the sampling↳
↳ strategy

import time
name = []
Accuracy = []
Specificity=[]
Sensetivity=[]
```

```

alg_auc=[]
T_time=[]
P_time=[]
Mcc=[]
br_s=[]
f1score=[]
for alg in RUSMLA:
    name.append(alg.__class__.__name__)

    #resampling
    x_Train_rus,y_Train_rus=RUS.fit_resample(x_train,y_train)

    #train the models
    start=time.time()
    alg.fit(x_Train_rus,y_Train_rus)
    end=time.time()

    #prediting the test set by each model
    startp=time.time()
    y_predicted=alg.predict(x_test)
    endp=time.time()

    #calculating accuracy, specificity, sensetivity, MCC, brier score,F1 Score
    Accuracy.append(accuracy_score(y_test,y_predicted))
    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))
    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))
    Mcc.append(matthews_corrcoef(y_test, y_predicted))
    br_s.append(brier_score_loss(y_test, y_predicted))
    f1score.append(f1_score(y_test, y_predicted))

    #computing ROC-AUC
    alg_probs= alg.predict_proba(x_test)
    alg_probs=alg_probs[:,1]
    alg_auc.append(roc_auc_score(y_test,alg_probs))

    #counting the running time of prediting and training mdoel
    T_time.append(end-start)
    P_time.append(endp-startp)

RUScomparison_NotTuned = pd.DataFrame({"name": name, "Accuracy":_
→Accuracy,"Sensetivity":Sensetivity,
                                     "Specificity":Specificity,"AUC":
→alg_auc,"Mcc":Mcc,"brier_score_loss":br_s,"F1Score":f1score,
                                     "Training Time [s]": T_time,"Prediction_
→Time [s]": P_time})
display(RUScomparison_NotTuned)

```

*#If we compared the below result with the results of Smotecomparison,
#we can notice the importance of the tuning as most of these measurements have
→ improved.*

	name	Accuracy	Sensetivity	Specificity	AUC	\
0	DecisionTreeClassifier	0.808050	0.792135	0.810184	0.801159	
1	LogisticRegression	0.842853	0.808989	0.847394	0.907632	
2	RandomForestClassifier	0.836610	0.890449	0.829391	0.922392	

	Mcc	brier_score_loss	F1Score	Training Time [s]	\
0	0.442796	0.191950	0.493870	0.070006	
1	0.503436	0.157147	0.548990	1.119425	
2	0.532751	0.163390	0.563055	1.056865	

	Prediction Time [s]
0	0.007898
1	0.005370
2	0.367377

5.4 Random Forest Feature Importance

The important features for Random Forest have been extracted based on the best balancing strategy that improved the performance of the Random Forest, which is the random undersampling

```
[56]: #important note: The model parameters have been adjusted based on the best  

      →balancing method.  

      rf=RandomForestClassifier(n_estimators=RUS_RF_n_estimators,  

                               max_depth= RUS_RF_max_depth,random_state=3)  

      rf.fit(rfx_train_rus,rfy_train_rus)  

      import matplotlib.pyplot as plt  

      import seaborn as sns  

      features=dff.loc[:, dff.columns != 'y'].columns  

      importances=rf.feature_importances_  

      indices=np.argsort(importances)  

      f = plt.figure()  

      f.set_figwidth(10)  

      f.set_figheight(20)  

      plt.title("Feature Importance")  


      plt.  

      →barh(range(len(indices)),importances[indices],color="b",align="center",height=  

      →8)  

      plt.yticks(range(len(indices)),[features[i] for i in indices])
```

```
plt.xlabel("Relative Importance (Gini importance)")  
plt.show()
```



```
[57]: import pandas as pd
df_1=pd.DataFrame({"Feature names":features,"Importances":rf.
    ↳feature_importances_})
df_1=df_1.sort_values(by="Importances",ascending=False)
df_1
```

```
[57]:
```

	Feature names	Importances
3	duration	0.319455
0	age	0.070855
1	balance	0.062890
2	day	0.057665
49	poutcome_success	0.050000
5	pdays	0.040041
4	campaign	0.028371
28	housing_no	0.021126
29	housing_yes	0.020490
6	previous	0.020401
34	contact_unknown	0.019647
50	poutcome_unknown	0.019590
42	month_mar	0.017013
32	contact_cellular	0.015839
35	month_apr	0.015358
45	month_oct	0.012146
41	month_jun	0.011714
47	poutcome_failure	0.010657
36	month_aug	0.010271
43	month_may	0.010256
24	education_tertiary	0.008878
20	marital_married	0.008740
38	month_feb	0.008499
40	month_jul	0.008367
21	marital_single	0.008220
23	education_secondary	0.008044
46	month_sep	0.007814
44	month_nov	0.007487
31	loan_yes	0.007337
16	job_technician	0.007183
8	job_blue-collar	0.007081
11	job_management	0.007063
30	loan_no	0.006662
7	job_admin.	0.006015
22	education_primary	0.005622
19	marital_divorced	0.005567
15	job_student	0.004826
39	month_jan	0.004659

33	contact_telephone	0.004551
12	job_retired	0.004491
14	job_services	0.004293
48	poutcome_other	0.003913
25	education_unknown	0.003635
13	job_self-employed	0.003431
37	month_dec	0.003419
17	job_unemployed	0.002842
9	job_entrepreneur	0.002418
10	job_housemaid	0.002268
27	default_yes	0.001156
26	default_no	0.000992
18	job_unknown	0.000741

```
[58]: #features with importance more than or equal the mean:
df_1[df_1["Importances"]>=df_1["Importances"].mean()]
```

```
[58]:
```

	Feature names	Importances
3	duration	0.319455
0	age	0.070855
1	balance	0.062890
2	day	0.057665
49	poutcome_success	0.050000
5	pdays	0.040041
4	campaign	0.028371
28	housing_no	0.021126
29	housing_yes	0.020490
6	previous	0.020401
34	contact_unknown	0.019647

```
[59]: #or we can just use SelectFromModel to extract the selected features
# SelectFromModel accepts a threshold parameter
#and will select the features whose importance (defined by the coefficients)
↳are above this threshold.
# the threshold is the mean by default
from sklearn.feature_selection import SelectFromModel
rf=ensemble.RandomForestClassifier(n_estimators=RUS_RF_n_estimators,
                                  max_depth=RUS_RF_max_depth,random_state=3)
rf.fit(rfx_train_rus,rfy_train_rus)

FS = SelectFromModel(rf)
FS.fit(rfx_train_rus, rfy_train_rus)
FS.get_support()
Important_Features= rfx_train_rus.columns[(FS.get_support())]
impresult=pd.DataFrame({"Important_Features": Important_Features})
display(impresult)
```

```

Important_Features
0          age
1         balance
2          day
3        duration
4        campaign
5         pdays
6        previous
7    housing_no
8    housing_yes
9    contact_unknown
10   poutcome_success

```

5.5 DT Feature Importance

```

[60]: #important note: we have to adjust the model parameters based on the best_
      ↪balancing method.


dt=DecisionTreeClassifier(max_depth=RUS_DT_max_depth,random_state=3)# rfgridRUS_
      ↪or rfgridSmote or Urfgrid
dt.fit(rfx_train_rus,rfy_train_rus) # (rfx_train_rus,rfy_train_rus) or_
      ↪(rfx_train_res,rfy_train_res) or (x_train,y_train)

import matplotlib.pyplot as plt
import seaborn as sns

features=dff.loc[:, dff.columns != 'y'].columns
importances=dt.feature_importances_
indices=np.argsort(importances)

f = plt.figure()
f.set_figwidth(10)
f.set_figheight(20)
plt.title("Feature Importance")
plt.
      ↪barh(range(len(indices)),importances[indices],color="b",align="center",height=.
      ↪8)
plt.yticks(range(len(indices)),[features[i] for i in indices])
plt.xlabel("Relative Importance (Gini importance)")
plt.show()

```

```
[61]: import pandas as pd
df_1=pd.DataFrame({"Feature names":features,"Importances":dt.
    ↳feature_importances_})
df_1=df_1.sort_values(by="Importances",ascending=False)
df_1
```

```
[61]:
```

	Feature names	Importances
3	duration	0.489557
49	poutcome_success	0.229189
29	housing_yes	0.057850
34	contact_unknown	0.051323
1	balance	0.032504
42	month_mar	0.019111
5	pdays	0.017468
0	age	0.016224
36	month_aug	0.016205
45	month_oct	0.011945
2	day	0.011352
35	month_apr	0.010435
46	month_sep	0.007474
20	marital_married	0.006382
43	month_may	0.003948
44	month_nov	0.003115
4	campaign	0.002938
10	job_housemaid	0.002374
28	housing_no	0.002042
48	poutcome_other	0.001320
22	education_primary	0.001251
11	job_management	0.001119
19	marital_divorced	0.001062
16	job_technician	0.001017
24	education_tertiary	0.000944
25	education_unknown	0.000930
26	default_no	0.000724
15	job_student	0.000197
39	month_jan	0.000000
6	previous	0.000000
7	job_admin.	0.000000
47	poutcome_failure	0.000000
8	job_blue-collar	0.000000
9	job_entrepreneur	0.000000
12	job_retired	0.000000
13	job_self-employed	0.000000
14	job_services	0.000000
41	month_jun	0.000000

40	month_jul	0.000000
38	month_feb	0.000000
27	default_yes	0.000000
37	month_dec	0.000000
17	job_unemployed	0.000000
18	job_unknown	0.000000
33	contact_telephone	0.000000
32	contact_cellular	0.000000
31	loan_yes	0.000000
30	loan_no	0.000000
21	marital_single	0.000000
23	education_secondary	0.000000
50	poutcome_unknown	0.000000

```
[62]: #features with importance more than or equal the mean:
df_1[df_1["Importances"]>=df_1["Importances"].mean()]
```

```
[62]:      Feature names  Importances
3      duration      0.489557
49  poutcome_success  0.229189
29      housing_yes   0.057850
34  contact_unknown   0.051323
1      balance       0.032504
```

5.6 Comparisons

5.6.1 Before SMOTE vs After SMOTE

```
[63]: for x in Smotecomparison.name:

    When=["Before smote","After smote"]
    smoteindex=Smotecomparison.loc[Smotecomparison['name'] == x].index.values[0]
    index=comparison.loc[comparison['name'] == x].index.values[0]

    Ac=[comparison.loc[comparison['name'] == x].Accuracy[index],
        Smotecomparison.loc[Smotecomparison['name'] == x].Accuracy[smoteindex]]
    Sen=[comparison.loc[comparison['name'] == x].Sensetivity[index],
        Smotecomparison.loc[Smotecomparison['name'] == x].
    ↪Sensetivity[smoteindex]]
    SP=[comparison.loc[comparison['name'] == x].Specificity[index],
        Smotecomparison.loc[Smotecomparison['name'] == x].Specificity[smoteindex]]
    ROC_AUC=[comparison.loc[comparison['name'] == x].AUC[index],
        Smotecomparison.loc[Smotecomparison['name'] == x].AUC[smoteindex]]
```

```


MCC2=[comparison.loc[comparison['name'] == x].MCC[index],
      Smotecomparison.loc[Smotecomparison['name'] == x].MCC[smoteindex]]
Brier=[comparison.loc[comparison['name'] == x].Brier_score[index],
       Smotecomparison.loc[Smotecomparison['name'] == x].
→Brier_score[smoteindex]]
F1=[comparison.loc[comparison['name'] == x].F1Score[index],
   Smotecomparison.loc[Smotecomparison['name'] == x].
→F1Score[smoteindex]]

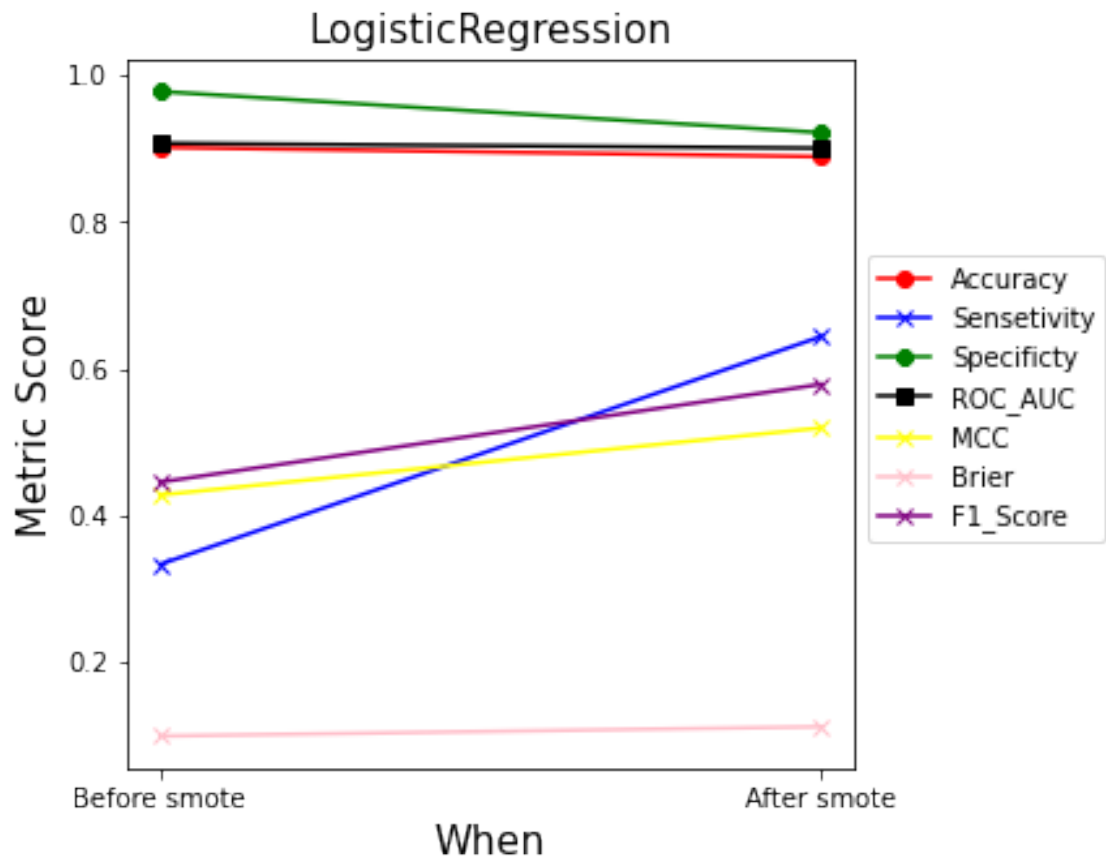
plt.figure(figsize=(5, 5))
plt.plot(When, Ac, color='red', marker='o')
plt.plot(When, Sen, color='blue', marker='x')
plt.plot(When, SP, color='green', marker='8')
plt.plot(When, ROC_AUC, color='black', marker='s')
plt.plot(When, MCC2, color='yellow', marker='x')
plt.plot(When, Brier, color='pink', marker='x')
plt.plot(When, F1, color='purple', marker='x')


plt.title(x, fontsize=15)
plt.xlabel('When', fontsize=15)
plt.ylabel('Metric Score', fontsize=15)
plt.legend(['Accuracy',
→'Sensetivity', 'Specificty', 'ROC_AUC', 'MCC', "Brier", "F1_Score"],
→fontsize=10, loc='lower left', bbox_to_anchor =(1, .3))

plt.show()

```







5.6.2 Before RUS vs After RUS

```
[64]: for x in RUScomparison.name:

    When=["Before RUS","After RUS"]
    RUSindex=RUScomparison.loc[RUScomparison['name'] == x].index.values[0]
    index=comparison.loc[comparison['name'] == x].index.values[0]

    Ac=[comparison.loc[comparison['name'] == x].Accuracy[index],
        RUScomparison.loc[RUScomparison['name'] == x].Accuracy[RUSindex]]
    Sen=[comparison.loc[comparison['name'] == x].Sensetivity[index],
        RUScomparison.loc[RUScomparison['name'] == x].Sensetivity[RUSindex]]
    SP=[comparison.loc[comparison['name'] == x].Specificity[index],
        RUScomparison.loc[RUScomparison['name'] == x].Specificity[RUSindex]]
    ROC_AUC=[comparison.loc[comparison['name'] == x].AUC[index],
        RUScomparison.loc[RUScomparison['name'] == x].AUC[RUSindex]]
    MCC2=[comparison.loc[comparison['name'] == x].MCC[index],
        RUScomparison.loc[RUScomparison['name'] == x].Mcc[RUSindex]]
```

```


Brier=[comparison.loc[comparison['name'] == x].Brier_score[index],
        RUScomparison.loc[RUScomparison['name'] == x].
↳brier_score_loss[RUSindex]]
F1=[comparison.loc[comparison['name'] == x].F1Score[index],
        RUScomparison.loc[RUScomparison['name'] == x].F1Score[RUSindex]]


plt.figure(figsize=(5, 5))
plt.plot(When, Ac, color='red', marker='o')
plt.plot(When, Sen, color='blue', marker='x')
plt.plot(When, SP, color='green', marker='8')
plt.plot(When, ROC_AUC, color='black', marker='s')
plt.plot(When, MCC2, color='yellow', marker='x')
plt.plot(When, Brier, color='pink', marker='x')
plt.plot(When, F1, color='purple', marker='x')

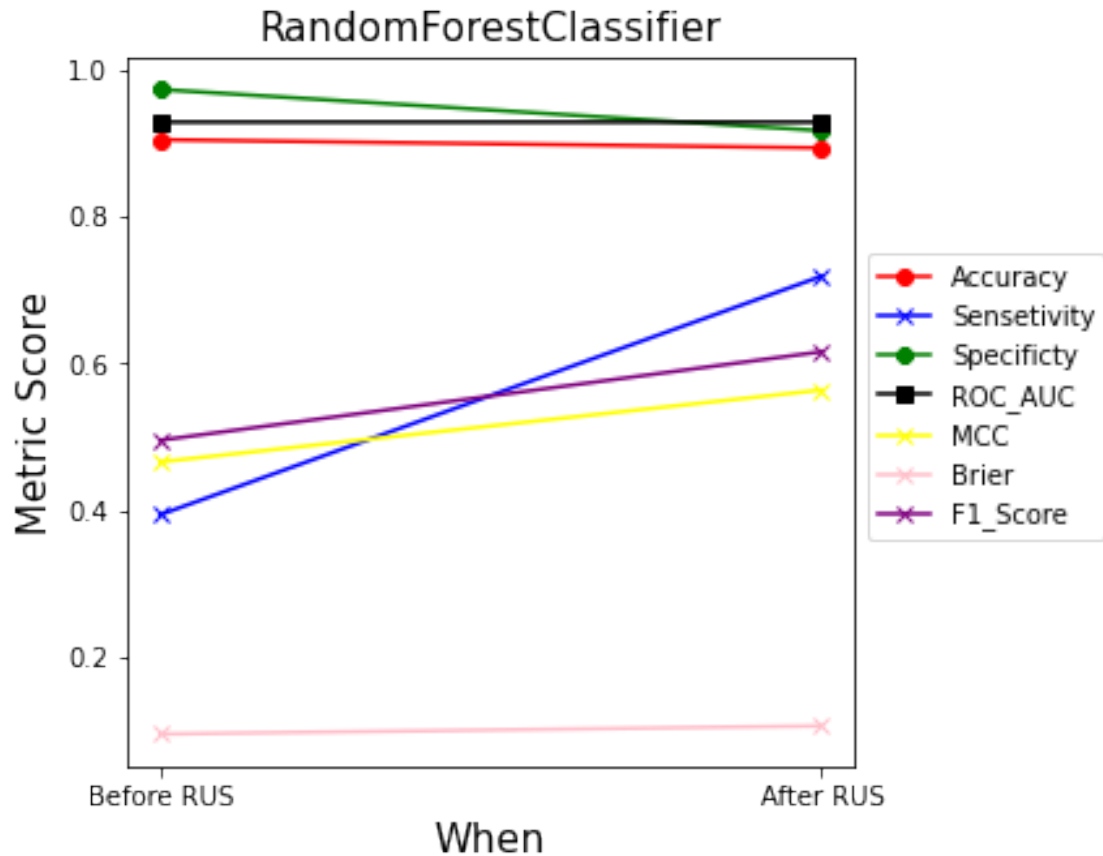
plt.title(x, fontsize=15)
plt.xlabel('When', fontsize=15)
plt.ylabel('Metric Score', fontsize=15)
plt.legend(['Accuracy',
↳'Sensetivity', 'Specificty', 'ROC_AUC', 'MCC', "Brier", "F1_Score", ],
↳
↳fontsize=10, loc='lower left', bbox_to_anchor =(1, .3))

plt.show()

```





5.6.3 SMOTE VS RUS

```
[65]: for x in Smotecomparison.name:

    When=["RUS", "SMOTE"]
    smoteindex=Smotecomparison.loc[Smotecomparison['name'] == x].index.values[0]
    RUSindex=RUScomparison.loc[RUScomparison['name'] == x].index.values[0]

    Ac=[RUScomparison.loc[RUScomparison['name'] == x].Accuracy[RUSindex],
        Smotecomparison.loc[Smotecomparison['name'] == x].Accuracy[smoteindex]]
    Sen=[RUScomparison.loc[RUScomparison['name'] == x].Sensitivity[RUSindex],
        Smotecomparison.loc[Smotecomparison['name'] == x].
    ↪Sensitivity[smoteindex]]
    SP=[RUScomparison.loc[RUScomparison['name'] == x].Specificity[RUSindex],
        Smotecomparison.loc[Smotecomparison['name'] == x].Specificity[smoteindex]]
    ROC_AUC=[RUScomparison.loc[RUScomparison['name'] == x].AUC[RUSindex],
        Smotecomparison.loc[Smotecomparison['name'] == x].AUC[smoteindex]]
```

```

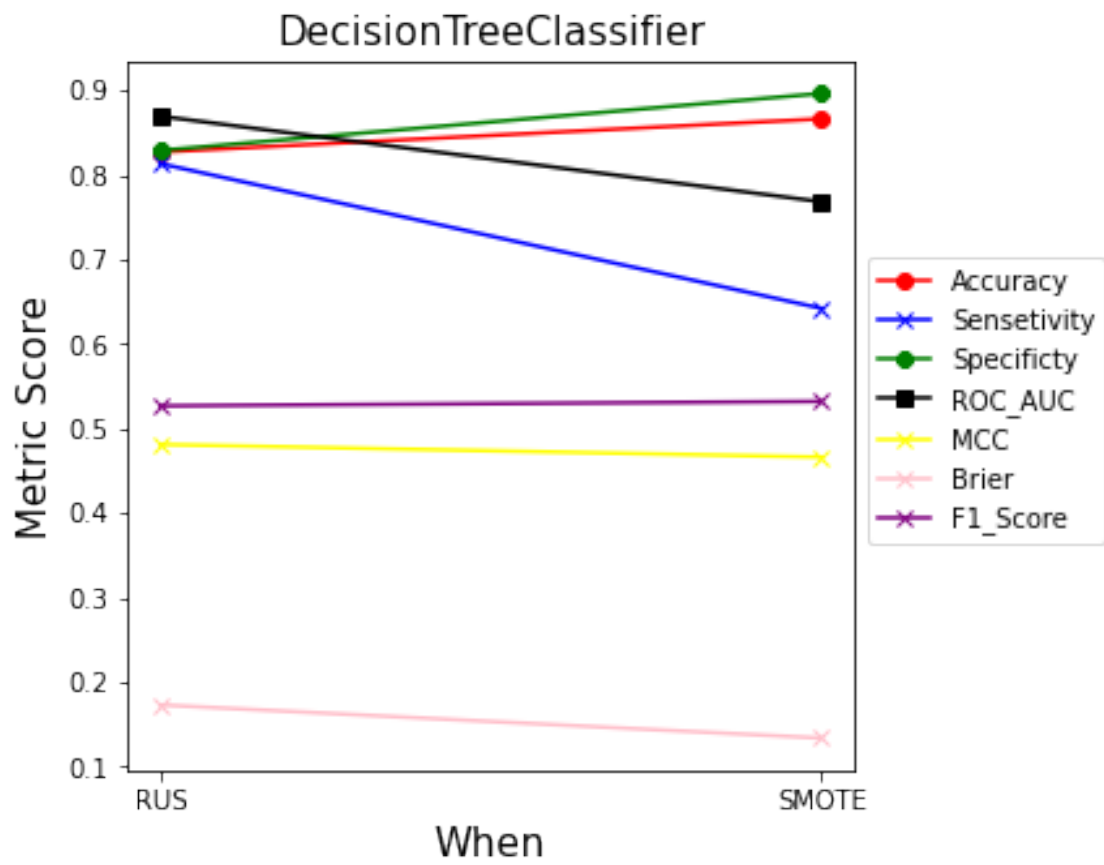
MCC2=[RUScomparison.loc[RUScomparison['name'] == x].Mcc[RUSindex],
      Smotecomparison.loc[Smotecomparison['name'] == x].MCC[smoteindex]]
Brier=[RUScomparison.loc[RUScomparison['name'] == x].
↳brier_score_loss[RUSindex],
      Smotecomparison.loc[Smotecomparison['name'] == x].
↳Brier_score[smoteindex]]
F1=[RUScomparison.loc[RUScomparison['name'] == x].F1Score[RUSindex],
    Smotecomparison.loc[Smotecomparison['name'] == x].
↳F1Score[smoteindex]]

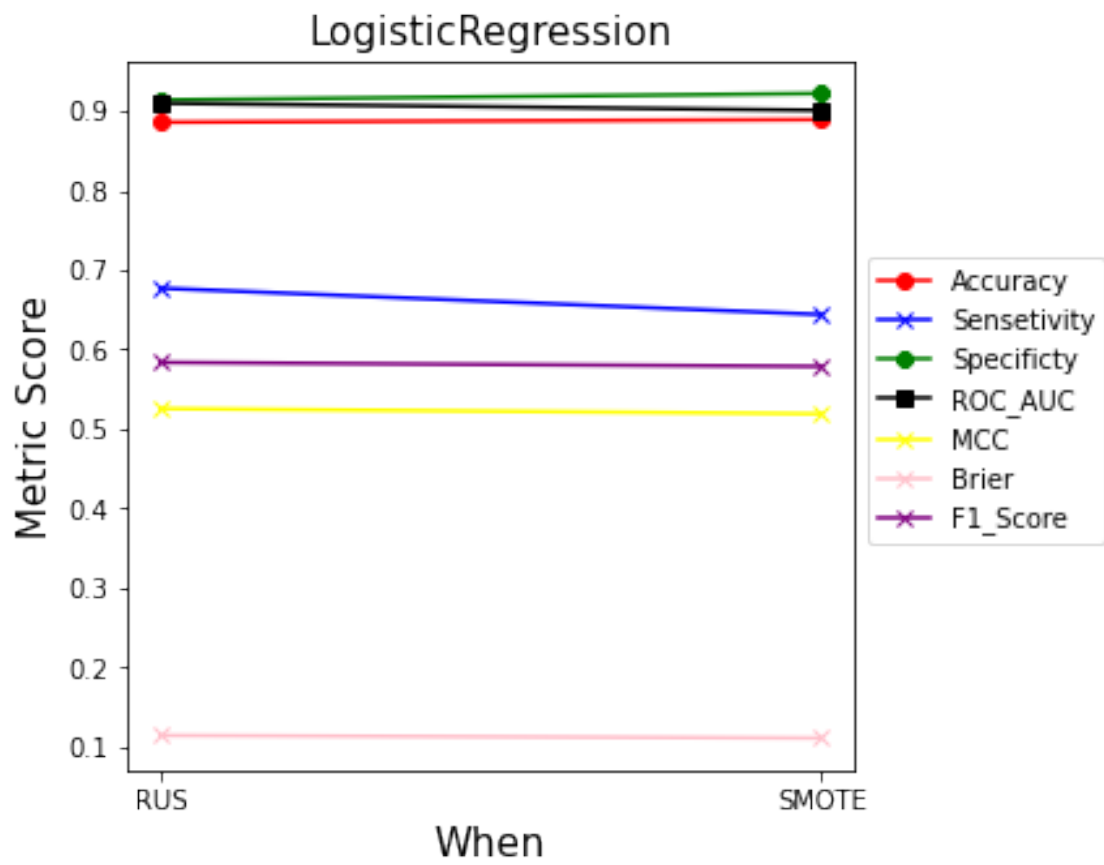
plt.figure(figsize=(5, 5))
plt.plot(When, Ac, color='red', marker='o')
plt.plot(When, Sen, color='blue', marker='x')
plt.plot(When, SP, color='green', marker='8')
plt.plot(When, ROC_AUC, color='black', marker='s')
plt.plot(When, MCC2, color='yellow', marker='x')
plt.plot(When, Brier, color='pink', marker='x')
plt.plot(When, F1, color='purple', marker='x')


plt.title(x, fontsize=15)
plt.xlabel('When', fontsize=15)
plt.ylabel('Metric Score', fontsize=15)
plt.legend(['Accuracy',
↳'Sensetivity', 'Specificty', 'ROC_AUC', 'MCC', "Brier", "F1_Score"],
↳
↳fontsize=10, loc='lower left', bbox_to_anchor =(1, .3))

plt.show()

```







5.6.4 RUS_Models Comparisons

```
[68]: df_T=RUScomparison.T
df_T

df_T.drop(index=df_T.index[0],
          axis=0,
          inplace=True)
df_T=df_T.rename(columns={0: 'DT',1:'LR',2:'RF'})

for x in df_T.index:

    When=["DT","LR","RF"]
    DTindex=df_T.loc[df_T.index == x]["DT"][0]
    LRindex=df_T.loc[df_T.index == x]["LR"][0]
    RFindex=df_T.loc[df_T.index == x]["LR"][0]

    if x=="Accuracy":
```

```

Ac=[DTindex,LRindex,RFindex]

elif x=="Sensetivity":
    Sen=[DTindex,LRindex,RFindex]

elif x=="Specificity":
    SP=[DTindex,LRindex,RFindex]

elif x=="AUC":
    ROC_AUC=[DTindex,LRindex,RFindex]

if x=="Mcc":
    MCC2=[DTindex,LRindex,RFindex]

if x=="brier_score_loss":
    Brier=[DTindex,LRindex,RFindex]


if x=="F1Score":
    F1=[DTindex,LRindex,RFindex]

plt.figure(figsize=(5, 5))
plt.plot(When, Ac, color='red', marker='o')
plt.plot(When, Sen, color='blue', marker='x')
plt.plot(When, SP, color='green', marker='8')
plt.plot(When, ROC_AUC, color='black', marker='s')
plt.plot(When, MCC2, color='orange', marker='x')
plt.plot(When, Brier, color='pink', marker='x')
plt.plot(When, F1, color='purple', marker='x')

plt.yticks(np.arange(min(Brier), max(Ac), 0.05))
plt.title("RUS_Models Comparesions", fontsize=15)
plt.xlabel('When', fontsize=15)
plt.ylabel('Metric Score', fontsize=15)
plt.legend(['Accuracy',
    ↳'Sensetivity', 'Specificty', 'ROC_AUC', 'MCC', "Brier", "F1_Score"],
    ↳fontsize=10, loc='lower left', bbox_to_anchor =(1, .3))

plt.show()

```

5.6.5 SMOTE_Models Comparasions

```
[69]: df_T=Smotecomparison.T
df_T

df_T.drop(index=df_T.index[0],
          axis=0,
          inplace=True)
df_T=df_T.rename(columns={0: 'DT',1:'LR',2:'RF'})

for x in df_T.index:

    When=["DT","LR","RF"]
    DTindex=df_T.loc[df_T.index == x]["DT"][0]
    LRindex=df_T.loc[df_T.index == x]["LR"][0]
    RFindex=df_T.loc[df_T.index == x]["LR"][0]

    if x=="Accuracy":
        Ac=[DTindex,LRindex,RFindex]
```

```

elif x=="Sensetivity":
    Sen=[DTindex,LRindex,RFindex]

elif x=="Specificity":
    SP=[DTindex,LRindex,RFindex]

elif x=="AUC":
    ROC_AUC=[DTindex,LRindex,RFindex]

if x=="Mcc":
    MCC2=[DTindex,LRindex,RFindex]

if x=="brier_score_loss":
    Brier=[DTindex,LRindex,RFindex]


if x=="F1Score":
    F1=[DTindex,LRindex,RFindex]

plt.figure(figsize=(5, 5))
plt.plot(When, Ac, color='red', marker='o')
plt.plot(When, Sen, color='blue', marker='x')
plt.plot(When, SP, color='green', marker='8')
plt.plot(When, ROC_AUC, color='black', marker='s')
plt.plot(When, MCC2, color='orange', marker='x')
plt.plot(When, Brier, color='pink', marker='x')
plt.plot(When, F1, color='purple', marker='x')

plt.title("SMOTE_Models Comparesions", fontsize=15)
plt.xlabel('When', fontsize=15)
plt.ylabel('Metric Score', fontsize=15)
plt.legend(['Accuracy',
    ↳'Sensetivity', 'Specificty', 'ROC_AUC', 'MCC', "Brier", "F1_Score"],
    ↳fontsize=10, loc='lower left', bbox_to_anchor =(1, .3))

plt.show()

```



6 Friedman Test to know how these models differ and which one is the best:

The Friedman Test is a non-parametric alternative to the Repeated Measures ANOVA. It is used to determine whether or not there is a statistically significant difference between three or more groups in which the same subjects show up in each group.

The procedure involves ranking each row (or block) together, then considering the values of ranks by columns. Applicable to complete block designs, it is thus a special case of the Durbin test.

The reason behind choosing Friedman test is that we want to compare multiple classifiers on multiple datasets (as we will apply the k fold cross validation (1n_splits=10, n_repeats=2)).

1. We will first create a data frame that has the result of AUC and MCC for each model on each fold.
2. we will apply the Friedman test to know if the models performance are statistically different or not .

-Please note that we already concluded that Random undersampling is the best practice to gain

better results and we decided that Random Forest will be the best model as its evaluation results were the best and it was stable on many applied scenarios.

```
[71]: from sklearn.metrics import roc_curve,   
      ↪ roc_auc_score, accuracy_score, recall_score, brier_score_loss  
from sklearn import tree, linear_model, ensemble  
from sklearn.metrics import mean_absolute_error  
from sklearn.model_selection import validation_curve  
from sklearn.model_selection import RepeatedKFold  
import numpy as np  
np.random.seed(3)  
import matplotlib.pyplot as plt  
plt.style.use("ggplot")
```

6.1 Matthews_corrcoef

```
[72]: #Matthews_corrcoef  
# prepare the cross-validation procedure  
#Due to the assumption that the test statistic has a chi squared distribution,   
      ↪ the p-value is only reliable for n > 10 and more than 6 repeated samples.  
cv = RepeatedKFold(n_splits=10, n_repeats=2, random_state=111)  
  
RUS_lgr_list_Test_MCC=[]  
RUS_RF_list_Test_MCC=[]  
RUS_DT_list_Test_MCC=[]  
  
#RUS_lgr_list_training_MCC=[]  
#RUS_RF_list_training_MCC=[]  
#RUS_DT_list_training_MCC=[]  
  
X=df.drop(["y"],axis="columns")  
y=df.y  
counter=0  
for train_index,test_index in cv.split(df):  
    counter+=1  
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]  
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]  
    log_reg=linear_model.  
    ↪ LogisticRegression(max_iter=RUS_lgr_max_iter,solver=RUS_lgr_solver,random_state=113,n_jobs=  
    DT=tree.DecisionTreeClassifier(max_depth=RUS_DT_max_depth,random_state=113)  
    RF=ensemble.RandomForestClassifier(n_estimators=RUS_RF_n_estimators,  
    ↪  
    ↪ max_depth=RUS_RF_max_depth,random_state=113)  
  
# Apply the tuned RUS to balance/resample the training set:
```

```

#Logistic Regression
↳
↳RUSlgr=RandomUnderSampler(sampling_strategy=lgr_RUS_sampling_strategy,random_state=11)
  lgrx_train_rus,lgry_train_rus=RUSlgr.fit_resample(X_train,Y_train)
#Decision Tree
↳
↳RUSdt=RandomUnderSampler(sampling_strategy=DT_RUS_sampling_strategy,random_state=11)
  dtx_train_rus,dtty_train_rus=RUSdt.fit_resample(X_train,Y_train)
#Random Forest
↳
↳RUSrf=RandomUnderSampler(sampling_strategy=RF_RUS_sampling_strategy,random_state=11)
  rfx_train_rus,rffy_train_rus=RUSrf.fit_resample(X_train,Y_train)

#Train the models on the resamples train sets
log_reg.fit(lgrx_train_rus,lgry_train_rus)#Logistic Regression
DT.fit(dtx_train_rus,dtty_train_rus)#Decision Tree
RF.fit(rfx_train_rus,rffy_train_rus)#Random Forest

#Predict the test set
log_reg_y_test_data_pred=log_reg.predict(X_test)
DT_y_test_data_pred=DT.predict(X_test)
RF_y_test_data_pred=RF.predict(X_test)

#Calculate MCC for each fold
↳
↳RUS_lgr_fold_testing_MCC=matthews_corrcoef(Y_test,log_reg_y_test_data_pred)#logistic
↳Regression
↳
↳RUS_DT_fold_testing_MCC=matthews_corrcoef(Y_test,DT_y_test_data_pred)#Decision
↳Tree
  RUS_RF_fold_testing_MCC=matthews_corrcoef(Y_test,RF_y_test_data_pred)#Random
↳Forest

  RUS_lgr_list_Test_MCC.append(RUS_lgr_fold_testing_MCC)#logistic Regression
  RUS_DT_list_Test_MCC.append(RUS_DT_fold_testing_MCC)#Decision Tree
  RUS_RF_list_Test_MCC.append(RUS_RF_fold_testing_MCC)#Random Forest

```

```

[73]: folds=[]
for x in range(counter):
    folds.append(x+1)
Results_MCC= pd.DataFrame({"Fold": folds,"DT":RUS_DT_list_Test_MCC,"LG":
↳RUS_lgr_list_Test_MCC,"RF":RUS_RF_list_Test_MCC})
display(Results_MCC)

```

	Fold	DT	LG	RF
0	1	0.488996	0.552656	0.590074
1	2	0.462728	0.522442	0.541534

2	3	0.471281	0.517129	0.571813
3	4	0.491908	0.520569	0.560187
4	5	0.462786	0.491912	0.554265
5	6	0.455930	0.532701	0.569286
6	7	0.477755	0.485883	0.564155
7	8	0.502080	0.502493	0.556888
8	9	0.447865	0.526346	0.554213
9	10	0.466944	0.519891	0.540321
10	11	0.470240	0.520860	0.562804
11	12	0.470980	0.541663	0.576400
12	13	0.488660	0.566835	0.583285
13	14	0.486593	0.504504	0.566506
14	15	0.444202	0.494670	0.548285
15	16	0.452667	0.522086	0.558773
16	17	0.490917	0.496121	0.563371
17	18	0.492884	0.505716	0.567325
18	19	0.488061	0.528497	0.555579
19	20	0.462488	0.476789	0.541600

```
[74]: from scipy import stats

#perform Friedman Test
stats.friedmanchisquare(Results_MCC["DT"], Results_MCC["LG"], Results_MCC["RF"])
# The Pvalue is less than 0.05 which means there is a significant difference
↳ between the three models
```

```
[74]: FriedmanchisquareResult(statistic=40.0, pvalue=2.0611536224385566e-09)
```

```
[ ]: #Interpret the results
#The Friedman Test uses the following null and alternative hypotheses:

#The null hypothesis (H0): There no difference between the models..

#The alternative hypothesis: (Ha): At least one model differs from the rest.

#Since this pvalue=2.0611536224385566e-09 which is less than 0.05, we reject
↳ the null hypothesis that tsays all the models are similar to each other.
#This returned a statistical difference, but now I would like to find out
↳ between which models the differences exist.

#The are a number of possibilities to perform posthoc-tests (an extension
↳ regarding the use of non-parametric tests :
#1 Perform the Nemenyi-test for all pairwise combinations; this is similar to
↳ the Tukey-test for ANOVA.
#2. Perform the Bonferroni-Dunn-test; in this setting one compares all values
↳ to a list of control values.
```

6.1.1 Perform the Nemenyi Test

Nemenyi Test: The Friedman Test is used to find whether there exists a significant difference between more than two groups. In such groups, the same subjects show up in each group. If the p-value of the Friedman test turns out to be statistically significant then we can conduct the Nemenyi test to find exactly which groups are different. This test is also known as Nemenyi post-hoc test.

The Friedman Test follows the below hypothesis:

1. The null hypothesis (H_0): There no difference between the models.
2. The alternative hypothesis: (H_a): At least one model differs from the rest.

```
[75]: pip install scikit-posthocs
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Collecting scikit-posthocs
  Downloading scikit_posthocs-0.7.0-py3-none-any.whl (38 kB)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-
packages (from scikit-posthocs) (3.6.2)
Requirement already satisfied: pandas>=0.20.0 in /usr/local/lib/python3.8/dist-
packages (from scikit-posthocs) (1.3.5)
Requirement already satisfied: seaborn in /usr/local/lib/python3.8/dist-packages
(from scikit-posthocs) (0.11.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages
(from scikit-posthocs) (1.21.6)
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages
(from scikit-posthocs) (1.7.3)
Requirement already satisfied: statsmodels in /usr/local/lib/python3.8/dist-
packages (from scikit-posthocs) (0.12.2)
Requirement already satisfied: python-dateutil>=2.7.3 in
/usr/local/lib/python3.8/dist-packages (from pandas>=0.20.0->scikit-posthocs)
(2.8.2)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-
packages (from pandas>=0.20.0->scikit-posthocs) (2022.6)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-
packages (from python-dateutil>=2.7.3->pandas>=0.20.0->scikit-posthocs) (1.15.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib->scikit-posthocs)
(1.0.6)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.8/dist-
packages (from matplotlib->scikit-posthocs) (7.1.2)
Requirement already satisfied: pyparsing>=2.2.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib->scikit-posthocs)
(3.0.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-
packages (from matplotlib->scikit-posthocs) (21.3)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.8/dist-packages (from matplotlib->scikit-posthocs)
```

(1.4.4)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib->scikit-posthocs) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.8/dist-packages (from matplotlib->scikit-posthocs) (4.38.0)

Requirement already satisfied: patsy>=0.5 in /usr/local/lib/python3.8/dist-packages (from statsmodels->scikit-posthocs) (0.5.3)

Installing collected packages: scikit-posthocs

Successfully installed scikit-posthocs-0.7.0

```
[76]: # Importing libraries
from scipy import stats
import scikit_posthocs as sp
import numpy as np
```

```
[78]: # Conduct the Nemenyi post-hoc test
data = np.array([Results_MCC["DT"], Results_MCC["LG"], Results_MCC["RF"]])
sp.posthoc_nemenyi_friedman(data.T)
```

```
[78]:          0          1          2
0  1.000000  0.004467  0.001000
1  0.004467  1.000000  0.004467
2  0.001000  0.004467  1.000000
```

```
[79]: block=[]
for x in range(len(Results_MCC["Fold"])):
    x+=1
    block.append(x)
FBlock=block+block+block

Results_MCC1=Results_MCC.drop(["Fold"],axis="columns")
longdf=pd.melt(Results_MCC1,var_name="models",value_name="MCC")
longdf["block"]=FBlock
longdf.head()

display(sp.
    ↳posthoc_nemenyi_friedman(longdf,y_col="MCC",group_col="models",block_col="block",melted=True))
```

	DT	LG	RF
DT	1.000000	0.004467	0.001000
LG	0.004467	1.000000	0.004467
RF	0.001000	0.004467	1.000000

```
[80]: #The Nemenyi post-hoc test produces the p-values for each pairwise comparison
↳ of means. These values are:
```



```
# P-value of Decision Tree MCC vs. Logistic Regression MCC: 0.004467
# P-value of Decision Tree MCC vs. Random Forest MCC: 0.001000
# P-value of Logistic Regression MCC vs. Random Forest MCC: 0.004467

# so all of them are different from each other
```

6.2 AUROC

```
[81]: from sklearn.metrics import roc_curve, roc_auc_score
# prepare the cross-validation procedure
#Due to the assumption that the test statistic has a chi squared distribution,
↳ the p-value is only reliable for n > 10 and more than 6 repeated samples.
cv = RepeatedKfold(n_splits=10, n_repeats=2, random_state=100)

RUS_lgr_list_Test_AUR=[]
RUS_RF_list_Test_AUR=[]
RUS_DT_list_Test_AUR=[]

X=dff.drop(["y"],axis="columns")
y=dff.y
counter=0
for train_index,test_index in cv.split(dff):
    counter+=1
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]
    Y_train,Y_test=y.iloc[train_index],y.iloc[test_index]

    log_reg=linear_model.
↳ LogisticRegression(max_iter=RUS_lgr_max_iter,solver=RUS_lgr_solver,random_state=113,n_jobs=
    DT=tree.DecisionTreeClassifier(max_depth=RUS_DT_max_depth,random_state=113)
    RF=ensemble.RandomForestClassifier(n_estimators=RUS_RF_n_estimators,
↳
↳ max_depth=RUS_RF_max_depth,random_state=113)

# Apply the tuned RUS to balance/resample the training set:
#Logistic Regression
↳
↳ RUSlgr=RandomUnderSampler(sampling_strategy=lgr_RUS_sampling_strategy,random_state=11)
    lgrx_train_rus,lgry_train_rus=RUSlgr.fit_resample(X_train,Y_train)
#Decision Tree
↳
↳ RUSdt=RandomUnderSampler(sampling_strategy=DT_RUS_sampling_strategy,random_state=11)
    dtx_train_rus,dty_train_rus=RUSdt.fit_resample(X_train,Y_train)
#Random Forest
↳
↳ RUSrf=RandomUnderSampler(sampling_strategy=RF_RUS_sampling_strategy,random_state=11)
```

```

rfx_train_rus,rfy_train_rus=RUSrf.fit_resample(X_train,Y_train)

#Train the models
log_reg.fit(lgrx_train_rus,lgry_train_rus)#Logistic Regression
DT.fit(dtx_train_rus,dt_y_train_rus)#Decision Tree
RF.fit(rfx_train_rus,rfy_train_rus)#Random Forest

#Predict the test set
log_reg_y_test_data_pred=log_reg.predict(X_test)
DT_y_test_data_pred=DT.predict(X_test)
RF_y_test_data_pred=RF.predict(X_test)

#to get the probabilities of the prediction for each instance in the test set

rf_probs= RF.predict_proba(X_test)
dt_probs= DT.predict_proba(X_test)
lgr_probs= log_reg.predict_proba(X_test)

#to keep the probabilities of positive outcomes
rf_probs=rf_probs[:,1]
dt_probs=dt_probs[:,1]
lgr_probs=lgr_probs[:,1]

#Calculate AUR for each fold

RUS_RF_fold_testing_AUR=roc_auc_score(Y_test,rf_probs)
RUS_DT_fold_testing_AUR=roc_auc_score(Y_test,dt_probs)
RUS_lgr_fold_testing_AUR=roc_auc_score(Y_test,lgr_probs)

RUS_lgr_list_Test_AUR.append(RUS_lgr_fold_testing_AUR)#logistic Regression
RUS_DT_list_Test_AUR.append(RUS_DT_fold_testing_AUR)#Decision Tree
RUS_RF_list_Test_AUR.append(RUS_RF_fold_testing_AUR)#Random Forest

```

```

[82]: folds=[]
for x in range(counter):
    folds.append(x+1)
Results_AUR = pd.DataFrame({"Fold": folds,"DT":RUS_DT_list_Test_AUR,"LG":
    ↪RUS_lgr_list_Test_AUR,"RF":RUS_RF_list_Test_AUR})
display(Results_AUR)

```

	Fold	DT	LG	RF
0	1	0.879790	0.899323	0.929742
1	2	0.878971	0.901120	0.922327
2	3	0.883752	0.910998	0.929956
3	4	0.888113	0.909914	0.932232
4	5	0.889347	0.913651	0.937251
5	6	0.881926	0.906999	0.925350

6	7	0.877982	0.908436	0.933061
7	8	0.883344	0.904250	0.928707
8	9	0.878467	0.905538	0.924177
9	10	0.884521	0.913012	0.936749
10	11	0.878411	0.905661	0.925791
11	12	0.887769	0.908910	0.928506
12	13	0.869416	0.902577	0.925565
13	14	0.886525	0.910353	0.935359
14	15	0.867020	0.906838	0.930517
15	16	0.881018	0.905133	0.931214
16	17	0.888122	0.905864	0.928019
17	18	0.881229	0.904847	0.929309
18	19	0.896555	0.919044	0.935738
19	20	0.867961	0.907986	0.926515

```
[83]: from scipy import stats

#perform Friedman Test
stats.friedmanchisquare(Results_AUR["DT"], Results_AUR["LG"], Results_AUR["RF"])
# The Pvalue is less than 0.05 which means there is a significant difference
→between the three models
```

```
[83]: FriedmanchisquareResult(statistic=40.0, pvalue=2.0611536224385566e-09)
```

6.2.1 Nemenyi Test

```
[84]: # Conduct the Nemenyi post-hoc test
data = np.array([Results_AUR["DT"], Results_AUR["LG"], Results_AUR["RF"]])
sp.posthoc_nemenyi_friedman(data.T)
```

```
[84]:
```

	0	1	2
0	1.000000	0.004467	0.001000
1	0.004467	1.000000	0.004467
2	0.001000	0.004467	1.000000

```
[ ]: #The Nemenyi post-hoc test produces the p-values for each pairwise compariso.
→These values are:

# P-value of Decision Tree MCC vs. Logistic Regression MCC: 0.004467
# P-value of Decision Tree MCC vs. Random Forest MCC: 0.001000
# P-value of Logistic Regression MCC vs. Random Forest MCC:: 0.004467

# So all of them are different from each other
```

7 Export the report

```
[ ]: #convert it to pdf
!sudo apt-get install texlive-xetex texlive-fonts-recommended
↪texlive-plain-generic
```

```
[86]: !jupyter nbconvert --to pdf /content/Final_Project_Code.ipynb
```

[illegible]

```
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Making directory ./Final_Project_Code_files
[NbConvertApp] Writing 451347 bytes to ./notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', './notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', './notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 1472896 bytes to /content/Final_Project_Code.pdf
```